

# Computer Vision: Graph Neural Networks

Siheng Chen 陈思衡

# Outline

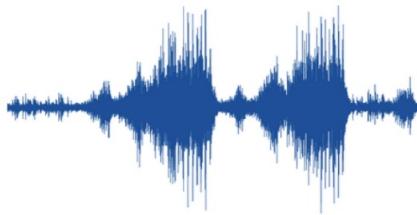
- Graph neural network basics
- GNN for visual physics
- GCN for multi-person behavior prediction

# Traditional Neural Networks

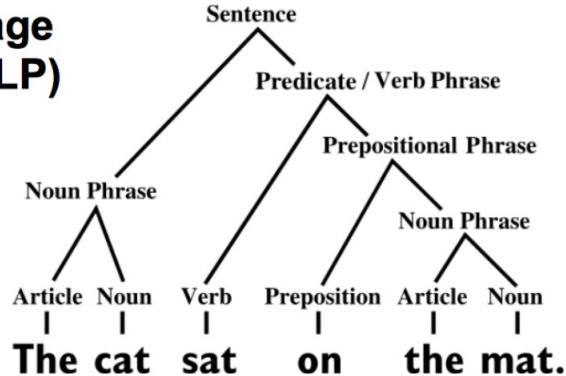
IMAGENET



Speech data

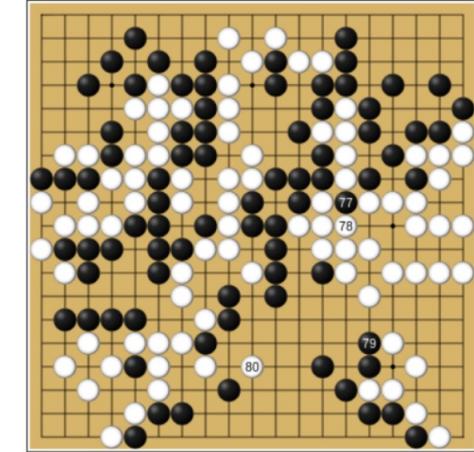


Natural language processing (NLP)



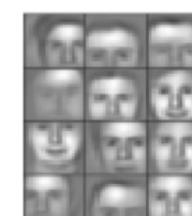
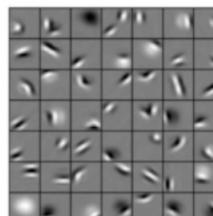
...

Grid games



Deep neural nets that exploit:

- translation equivariance (weight sharing)
- hierarchical compositionality



# Graph-Structured Data

A lot of real-world data does not “live” on grids

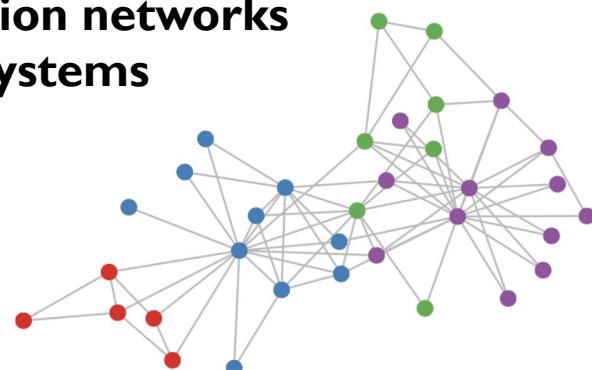
Standard CNN and RNN architectures  
don't work on these data

**Social networks**

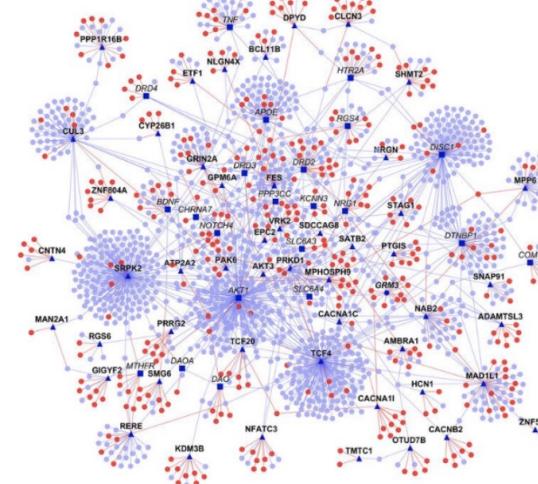
**Citation networks**

**Communication networks**

**Multi-agent systems**



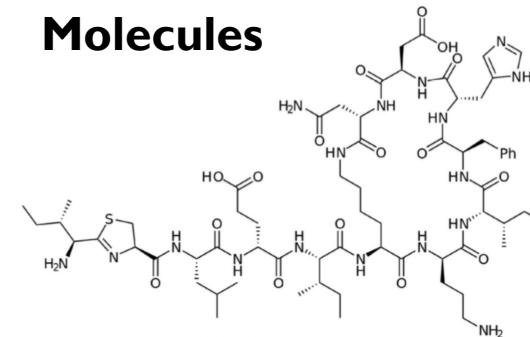
**Proteins interaction networks**



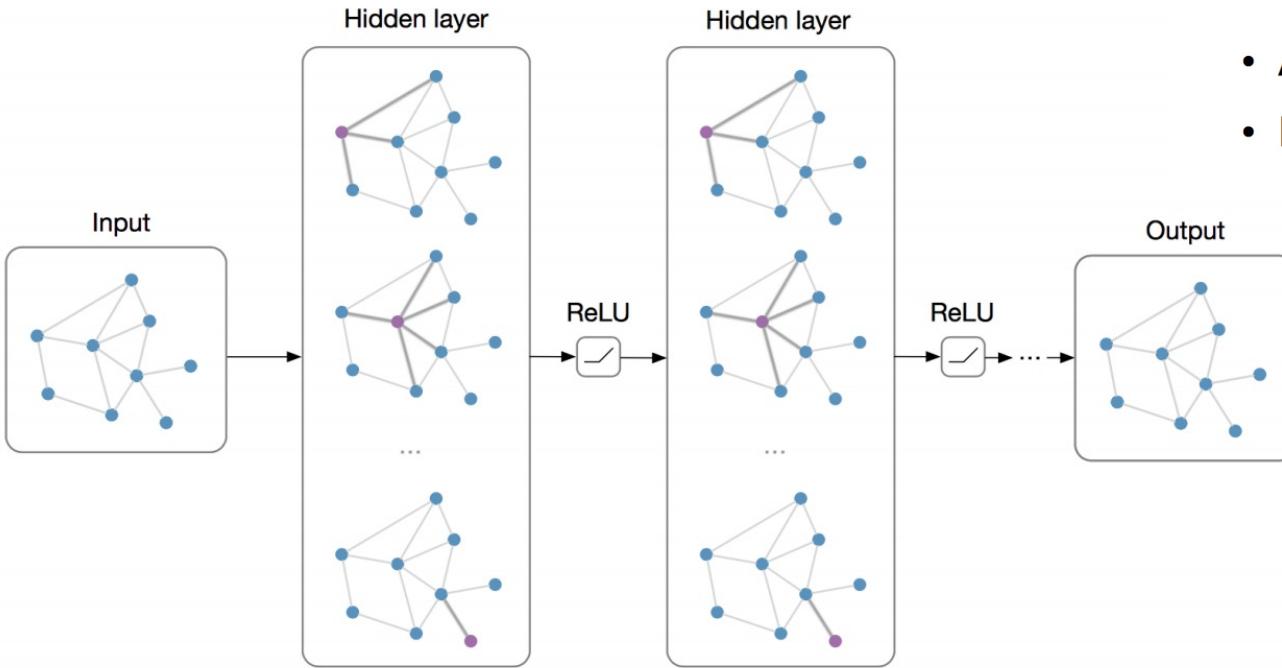
**Knowledge graphs**



**Molecules**



# Graph Neural Networks (GNNs)



**Notation:**  $\mathcal{G} = (\mathbf{A}, \mathbf{X})$

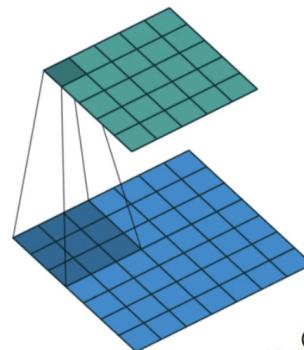
- Adjacency matrix  $\mathbf{A} \in \mathbb{R}^{N \times N}$
- Feature matrix  $\mathbf{X} \in \mathbb{R}^{N \times F}$

**Main Idea:** Pass messages between pairs of nodes and agglomerate

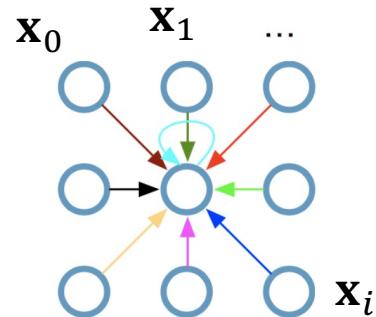
**Alternative Interpretation:** Pass messages between nodes to refine node  
(and possibly edge) representations

# Recap: Convolutional Neural Networks (CNNs) on Grids

## Single CNN layer with $3 \times 3$ filter



(Animation by  
Vincent Dumoulin)



$\mathbf{x}_i \in \mathbb{R}^F$  are (hidden layer) activations of a pixel/node

## Update for a single pixel:

- Transform messages individually  $\mathbf{W}_i \mathbf{x}_i$
- Add everything up  $\sum_i \mathbf{W}_i \mathbf{x}_i$

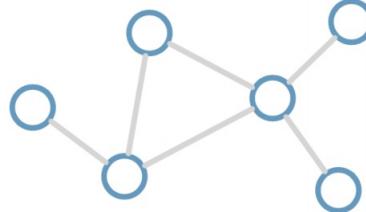
## Full update:

$$\bullet \quad \mathbf{x}_4^{(l+1)} = \sigma(\mathbf{W}_0^l \mathbf{x}_0^l + \mathbf{W}_1^l \mathbf{x}_1^l + \dots + \mathbf{W}_8^l \mathbf{x}_8^l)$$

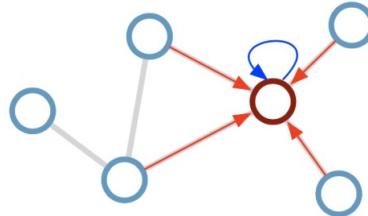
# Graph Convolutional Networks (GCNs)

Kipf & Welling (ICLR 2017), related previous works by Duvenaud et al. (NIPS 2015) and Li et al. (ICLR 2016)

Consider this undirected graph:



Calculate update for node in red:



## Desirable properties:

- Weight sharing over all locations
- Invariance to permutations
- Linear complexity  $O(E)$

**Update rule:**

$$\mathbf{x}_i^{(l+1)} = \sigma \left( \mathbf{W}_0^l \mathbf{x}_i^l + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{W}_1^l \mathbf{x}_j^{(l)} \right)$$

$\mathcal{N}_i$ : neighbor indices

$c_{ij}$ : norm. constant  
(fixed/trainable)

Matrix form



$$\mathbf{X}^{(l+1)} = \sigma(\tilde{\mathbf{A}} \mathbf{X}^{(l)} \mathbf{W})$$

$\mathbf{W} \in \mathbb{R}^{d_{(l+1)} \times d_{(l)}}$ : Weight matrix

$\tilde{\mathbf{A}} \in \mathbb{R}^{n \times n}$ : Normalized adjacency matrix

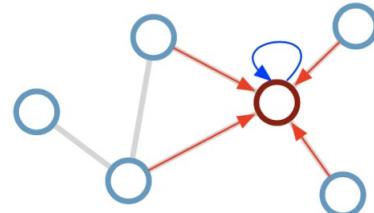
$\mathbf{X}^{(l)} \in \mathbb{R}^{n \times d_{(l)}}$ : Node feature vector in layer  $l$

# Graph Convolutional Networks (GCNs)

Kipf & Welling (ICLR 2017), related previous works by Duvenaud et al. (NIPS 2015) and Li et al. (ICLR 2016)

$$\mathbf{X}^{(l+1)} = \sigma(\tilde{\mathbf{A}}\mathbf{X}^{(l)}\mathbf{W})$$

feature aggregation



$$\mathbf{X}_{\text{out}} = \tilde{\mathbf{A}}\mathbf{X}_{\text{in}}$$

**Properties for feature aggregation :**

- Penalizes the feature difference between neighboring nodes
- low-pass filter

Interpreted as one gradient descent step for the  
**Laplacian smoothing problem :**

Make graph smoother

$$\arg \min_{\mathbf{X} \in \mathbb{R}^{n \times d}} \mathcal{L}_1(\mathbf{X}) := \frac{1}{2} \text{tr} \left( \mathbf{X}^\top (\mathbf{I} - \tilde{\mathbf{A}}) \mathbf{X} \right) = \frac{1}{2} \sum_{(v_i, v_j) \in \mathcal{E}} \left\| \frac{\mathbf{X}_i}{\sqrt{d_i + 1}} - \frac{\mathbf{X}_j}{\sqrt{d_j + 1}} \right\|_2^2$$

**One gradient descent step :**

$$\mathbf{X}_{\text{out}} = \mathbf{X}_{\text{in}} - (\mathbf{I} - \tilde{\mathbf{A}})\mathbf{X}_{\text{in}} = \tilde{\mathbf{A}}\mathbf{X}_{\text{in}}$$

# Graph Convolutional Networks (GCNs)

Kipf & Welling (ICLR 2017), related previous works by Duvenaud et al. (NIPS 2015) and Li et al. (ICLR 2016)

GCN Make graph smoother

Citation network: Node feature visualization with t-SNE



2 layers



5 layers



7 layers

2 layers: each cluster gets smoother

7 layers: the whole graph gets smoother (over smoothing)

# Graph Neural Networks with Attention (GAT)

Monti et al. (CVPR 2017), Hoshen (NIPS 2017), Veličković et al. (ICLR 2018)

GCN

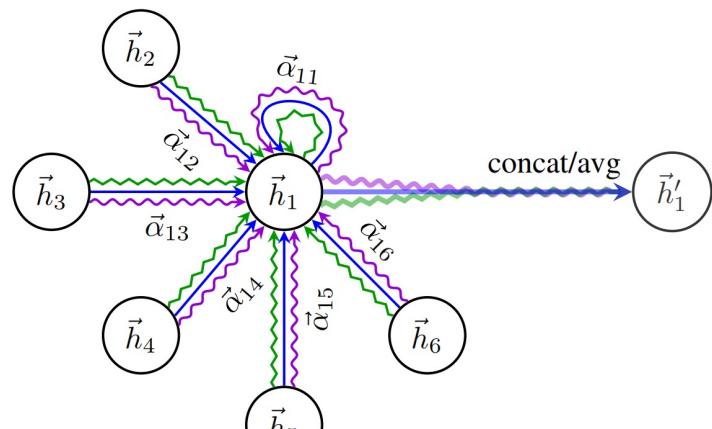
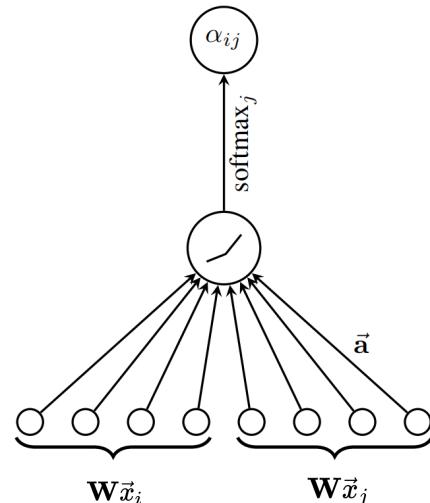
$$\mathbf{X}^{(l+1)} = \sigma(\tilde{\mathbf{A}}\mathbf{X}^{(l)}\mathbf{W})$$

General form of GAT

$$\mathbf{X}^{(l+1)} = \sigma((\alpha \circ \tilde{\mathbf{A}})\mathbf{X}^{(l)}\mathbf{W})$$

$\alpha$ : trainable attention matrix

$\tilde{\mathbf{A}}$ : prior adjacency mask



[Figure from Veličković et al. (ICLR 2018)]

$$\alpha_{ij} = \frac{\exp\left(\text{LeakyReLU}\left(\vec{a}^T [\mathbf{w}\vec{x}_i \| \mathbf{w}\vec{x}_j]\right)\right)}{\sum_{k \in \mathcal{N}_i} \exp\left(\text{LeakyReLU}\left(\vec{a}^T [\mathbf{w}\vec{x}_i \| \mathbf{w}\vec{x}_k]\right)\right)}$$

$$\vec{x}'_i = \sigma\left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{w}^k \vec{x}_j\right)$$

**Main idea: trainable adjacency coefficient (attention on neighbors)**

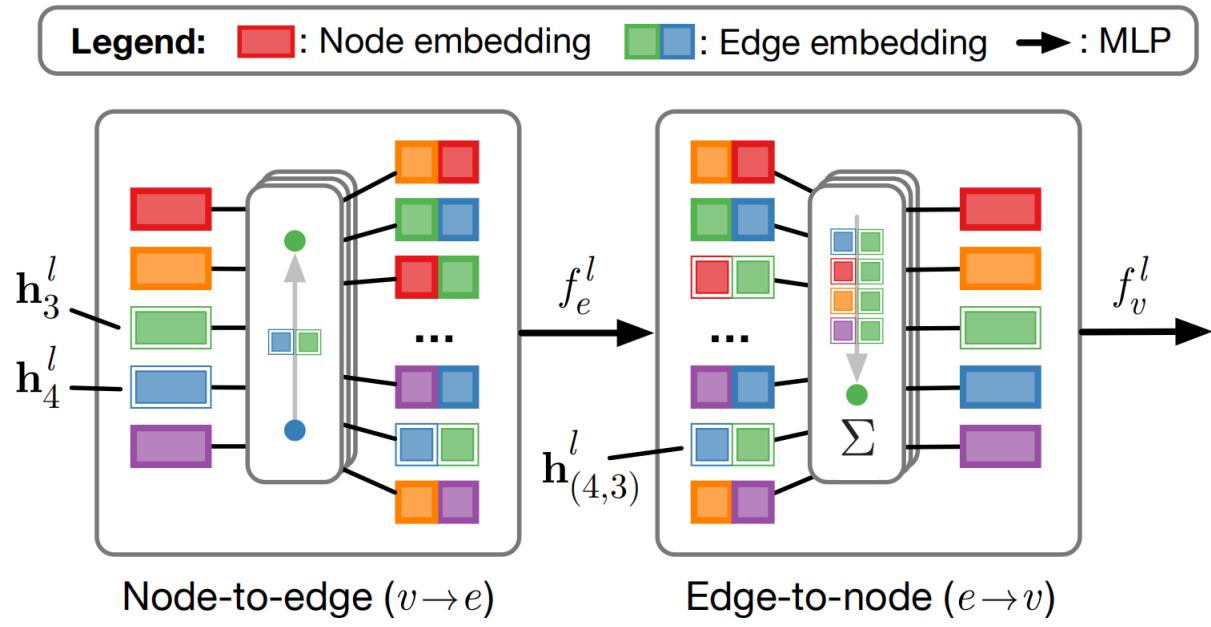
- Attention coefficient  $\alpha_{ij}$  from node  $i$  to  $j$
- Averaged multi-head-attention aggregation

**Properties for GAT :**

- Higher model capacity than GCN
- Efficient Computationally  
(parallelized across edges and nodes)
- Can be more difficult to optimize

# GNNs with Edge Embeddings

Battaglia et al. (NIPS 2016), Gilmer et al. (ICML 2017), Kipf et al. (ICML 2018)



Formally:

$$v \rightarrow e : \quad \mathbf{h}_{(i,j)}^l = f_e^l([\mathbf{h}_i^l, \mathbf{h}_j^l, \mathbf{x}_{(i,j)}])$$
$$e \rightarrow v : \quad \mathbf{h}_j^{l+1} = f_v^l([\sum_{i \in \mathcal{N}_j} \mathbf{h}_{(i,j)}^l, \mathbf{x}_j])$$

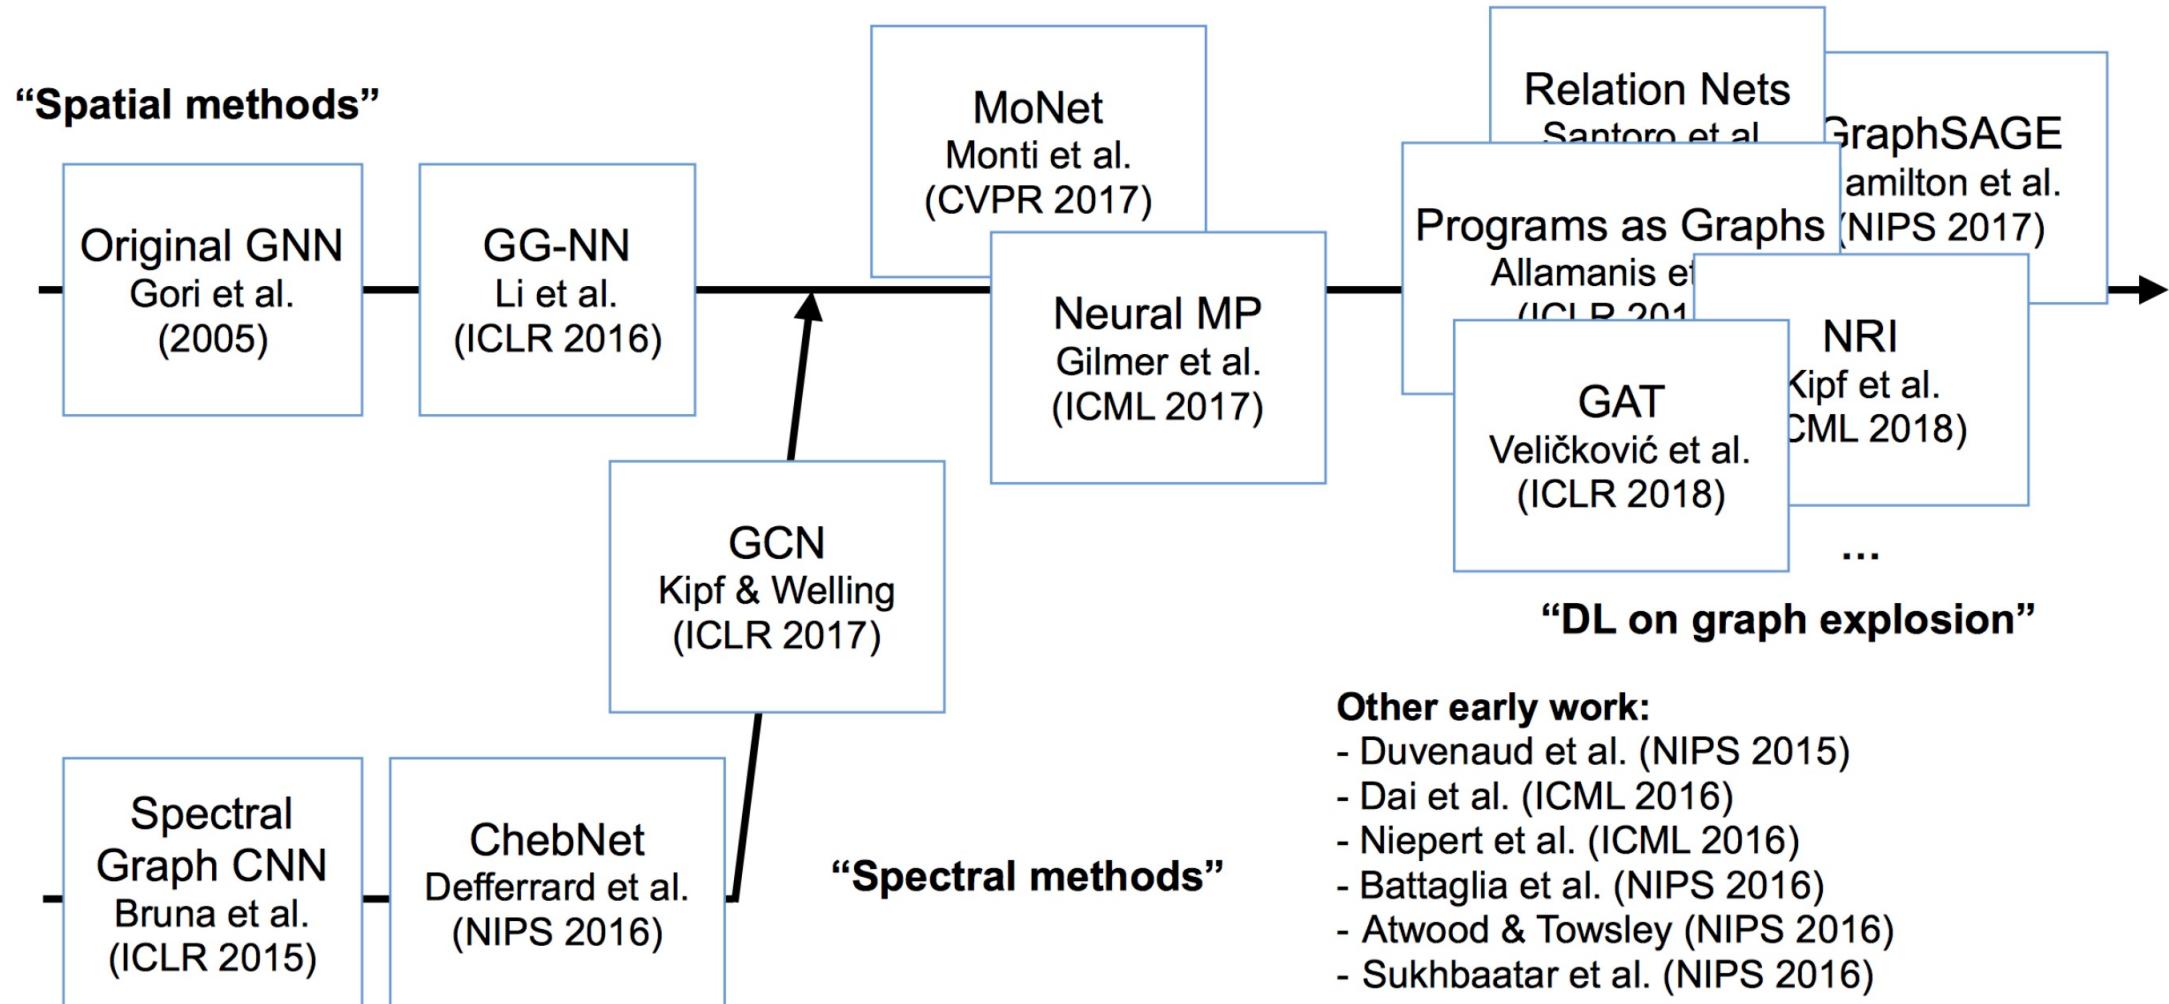
## Pros:

- Support edge features
- More expressive
- A general form of message passing

## Cons:

- Need to store intermediate edge-based activations
- In practice limited to small graphs

# A Brief History of Graph Neural Nets

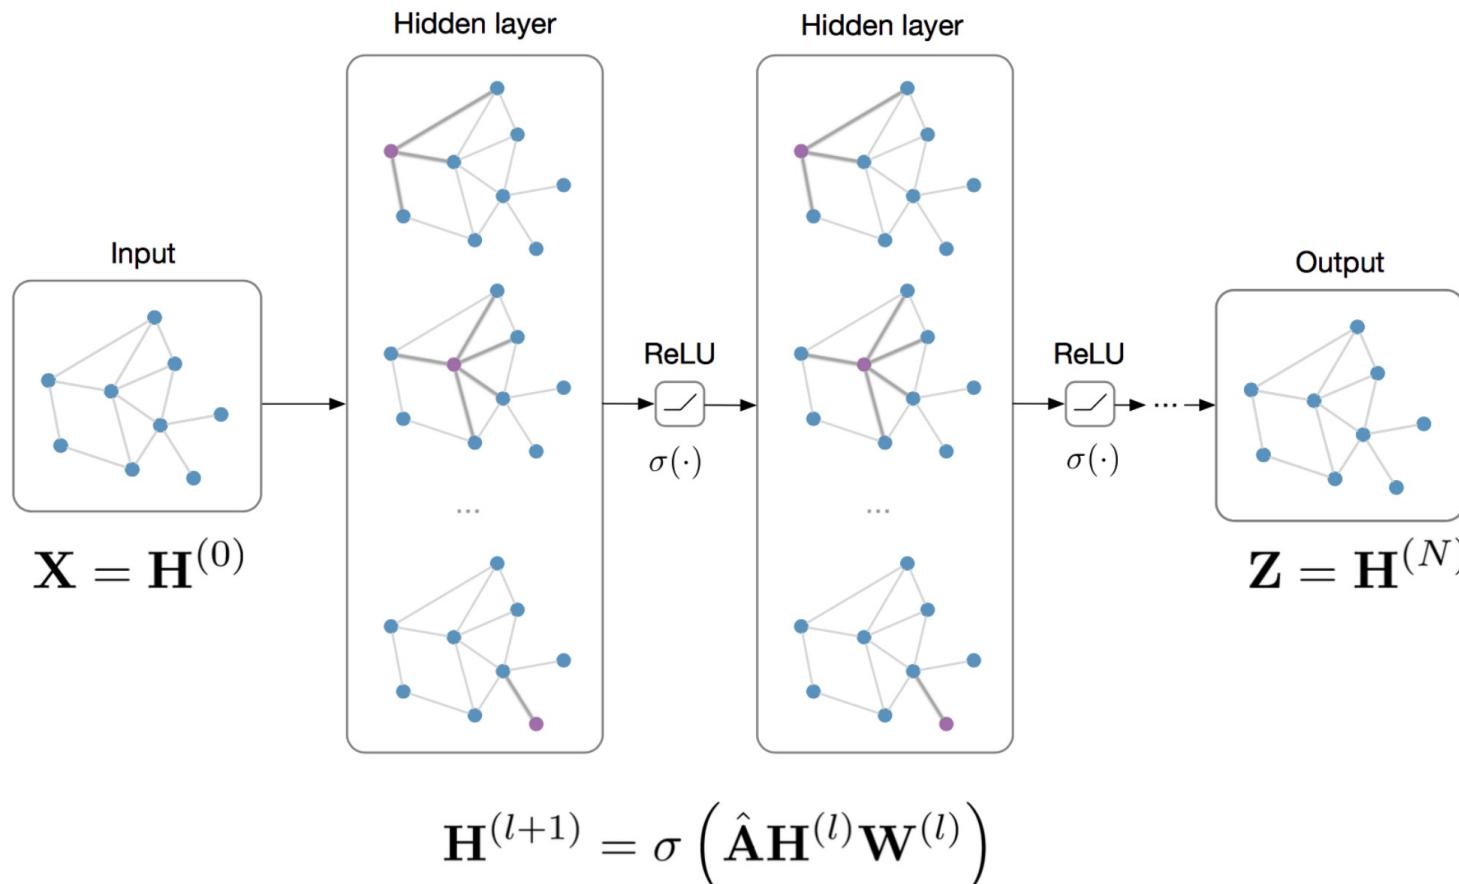


(slide inspired by Alexander Gaunt's talk on GNNs)

How do we use GNN / GCN for real problems?

# Classification and Link Prediction with GNNs / GCNs

Input: Feature matrix  $\mathbf{X} \in \mathbb{R}^{N \times E}$ , preprocessed adjacency matrix  $\hat{\mathbf{A}}$



**Node classification:**

$$\text{softmax}(\mathbf{z}_n)$$

e.g. Kipf & Welling (ICLR 2017)

**Graph classification:**

$$\text{softmax}(\sum_n \mathbf{z}_n)$$

e.g. Duvenaud et al. (NIPS 2015)

**Link prediction:**

$$p(A_{ij}) = \sigma(\mathbf{z}_i^T \mathbf{z}_j)$$

Kipf & Welling (NIPS BDL 2016)

“Graph Auto-Encoders”

# Semi-supervised Classification on Graphs

## Setting:

Some nodes are labeled (black circle)

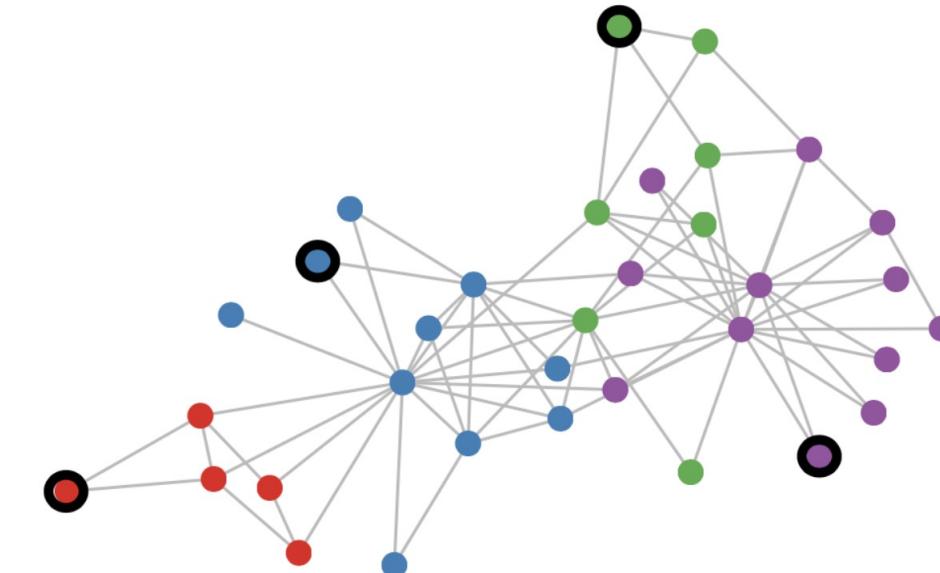
All other nodes are unlabeled

## Task:

Predict node label of unlabeled nodes

Evaluate loss on labeled nodes only:

$$\mathcal{L} = - \sum_{l \in \mathcal{Y}_L} \sum_{f=1}^F Y_{lf} \ln Z_{lf}$$



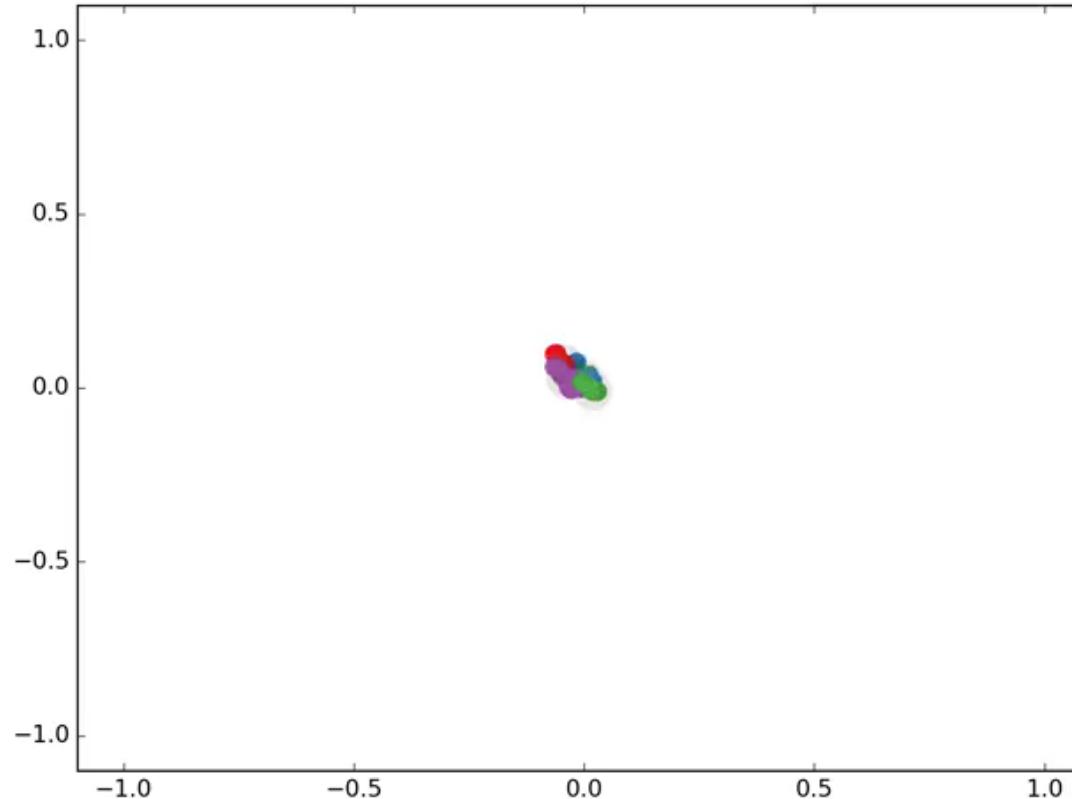
$\mathcal{Y}_L$  set of labeled node indices

$\mathbf{Y}$  label matrix

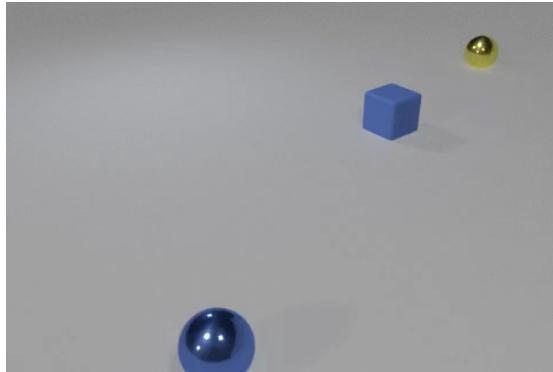
$\mathbf{Z}$  GCN output (after softmax)

# Semi-supervised Classification on Graphs

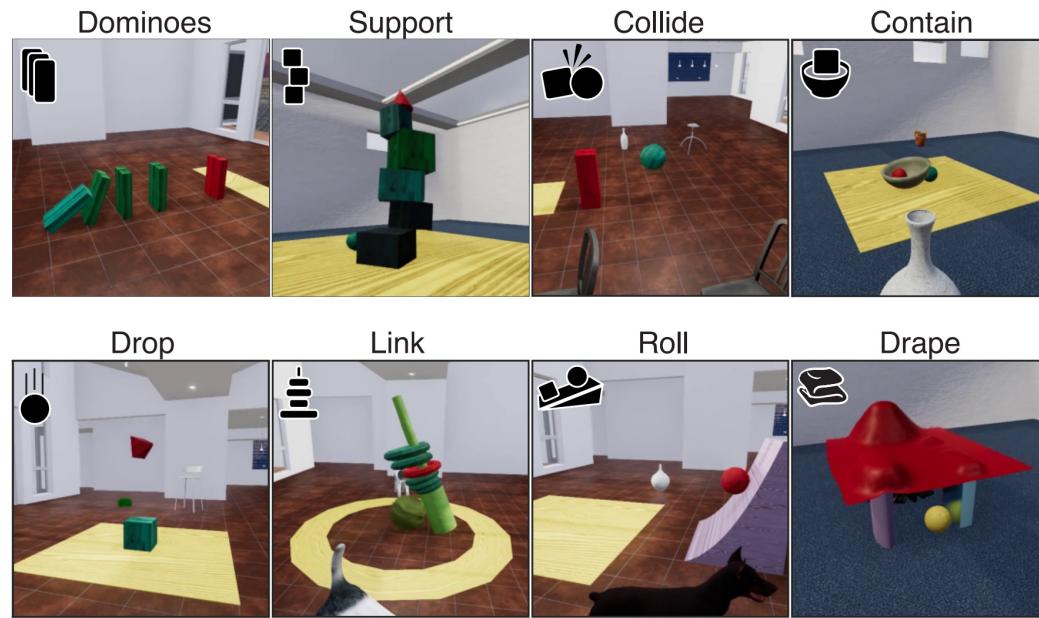
- Latent space dynamics for 300 training iterations
- Single label per class
- Labeled nodes are highlighted.



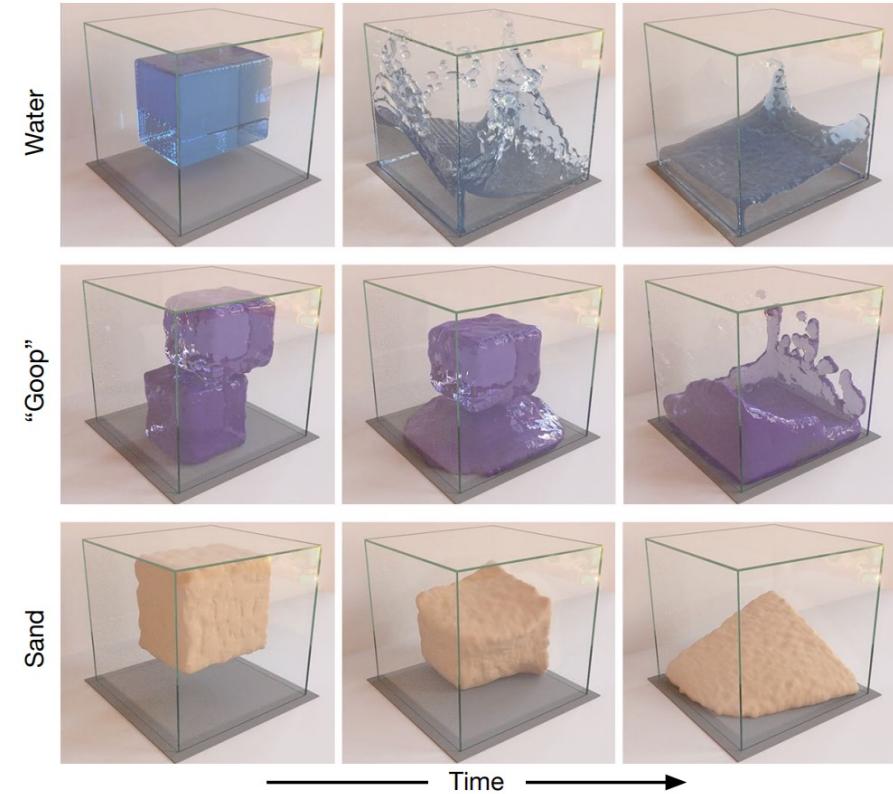
# GNN for physical model



Chen et al, Comphy dataset



Bear, Elias et al, Physion dataset



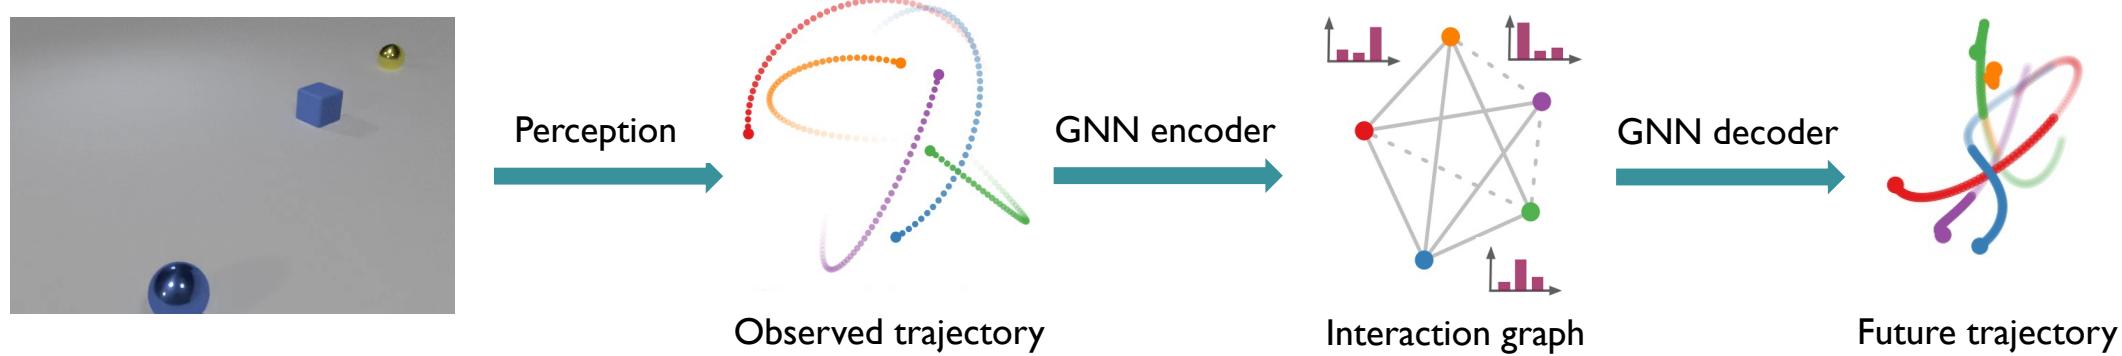
Alvaro et al, Learning to Simulate Complex Physics with Graph Networks

- Scene: dense or sparse physical objects
- Model: GNN to capture objects interaction

# GNN for relation inference

- What's the relation between objects? (same charged/opposite charged/no charge)
- The future trajectory of objects?

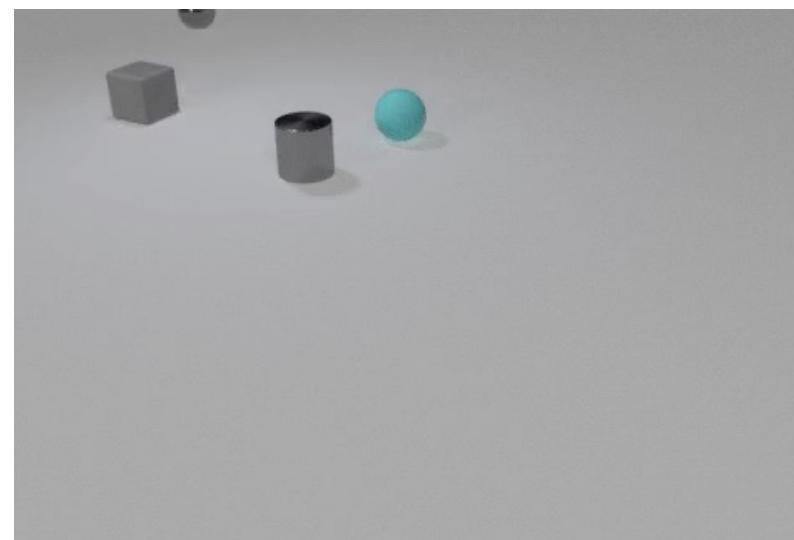
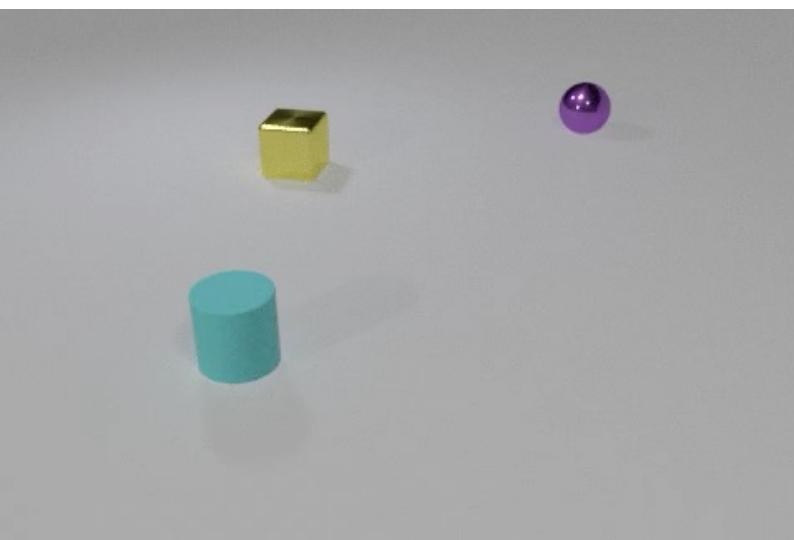
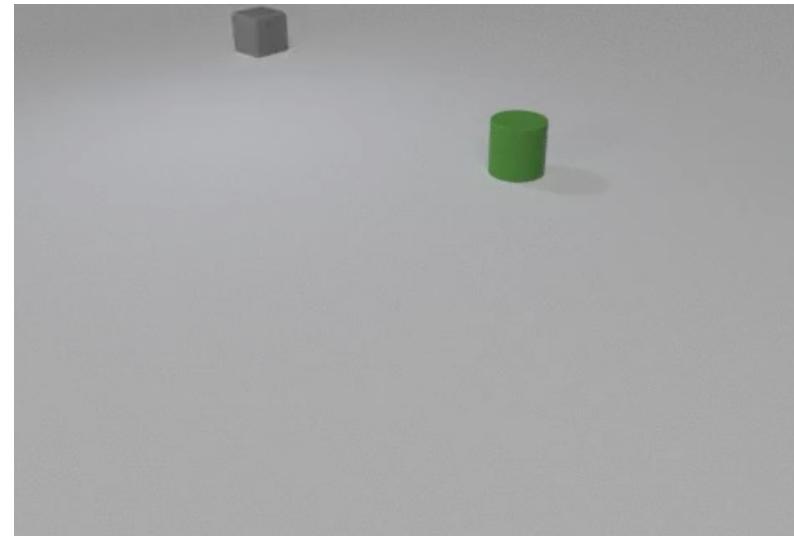
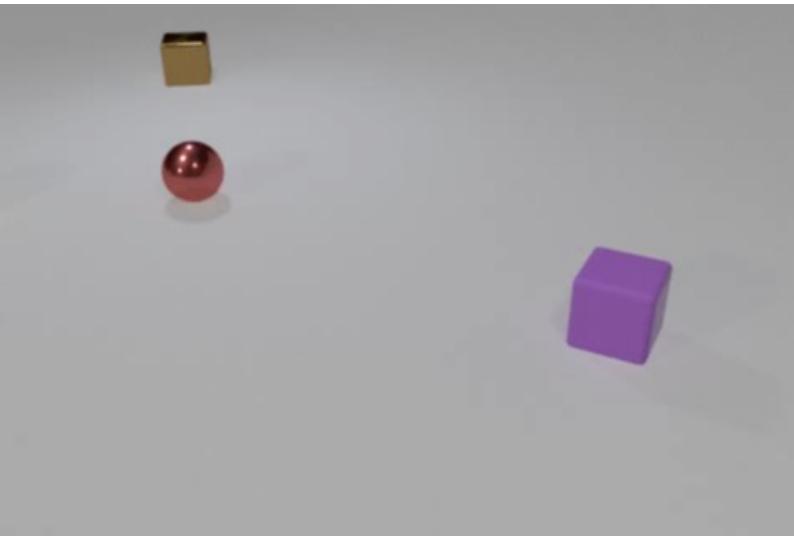
## Video



## Task framework:

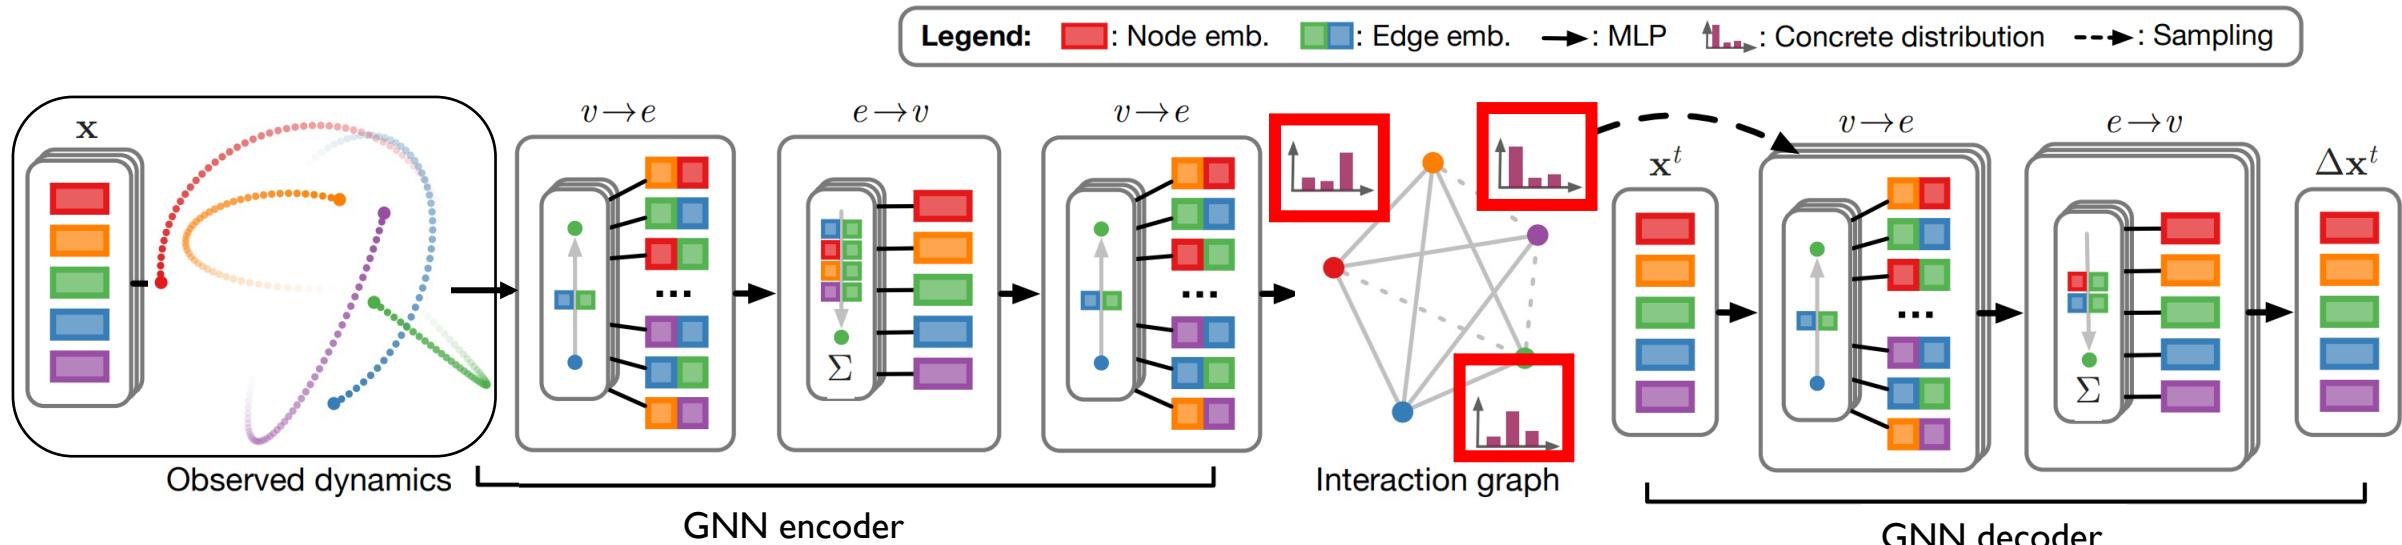
- Given video observation
- Video → trajectory
- Infer pair-wise relation of objects
- Predict future trajectory based on history trajectory and inferred relation

# GNN for relation inference



# GNN for relation inference

Kipf et al. Neural relational inference for interacting systems. ICML2018



$$\mathbf{X} \in \mathbb{R}^{T \times N \times d}$$

**T:** time steps

**N:** num of objects

**d:** motion feature dimension

$$\mathbf{h}_j^1 = f_{\text{emb}}(\mathbf{x}_j)$$

$$v \rightarrow e : \quad \mathbf{h}_{(i,j)}^1 = f_e^1([\mathbf{h}_i^1, \mathbf{h}_j^1])$$

$$e \rightarrow v : \quad \mathbf{h}_j^2 = f_v^1(\sum_{i \neq j} \mathbf{h}_{(i,j)}^1)$$

$$v \rightarrow e : \quad \mathbf{h}_{(i,j)}^2 = f_e^2([\mathbf{h}_i^2, \mathbf{h}_j^2])$$

$$v \rightarrow e : \quad \tilde{\mathbf{h}}_{(i,j)}^t = \sum_k z_{ij,k} \tilde{f}_e^k(\mathbf{x}_i^t, \mathbf{x}_j^t)$$

$$e \rightarrow v : \quad \boldsymbol{\mu}_j^{t+1} = \mathbf{x}_j^t + \tilde{f}_v(\sum_{i \neq j} \tilde{\mathbf{h}}_{(i,j)}^t)$$

$$p(\mathbf{x}_j^{t+1} | \mathbf{x}^t, \mathbf{z}) = \mathcal{N}(\boldsymbol{\mu}_j^{t+1}, \sigma^2 \mathbf{I})$$

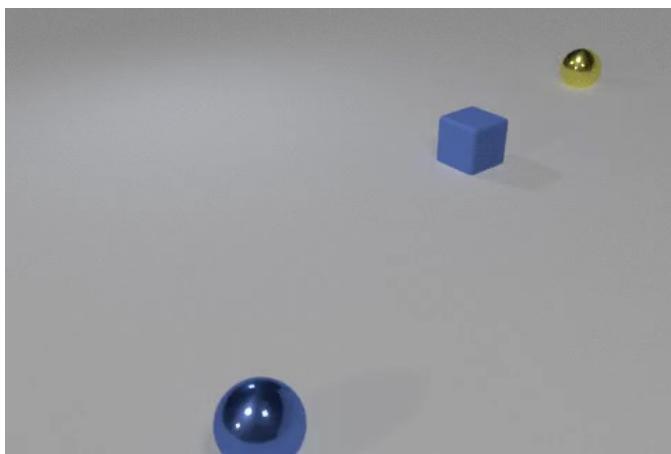
- GNN encodes edge embeddings for relation inference (can be supervised or unsupervised)
- **Different relationship types** correspond to different interaction modules in the GNN decoder
- GNN decodes node embeddings for trajectory prediction (supervised)

# GNN for relation inference

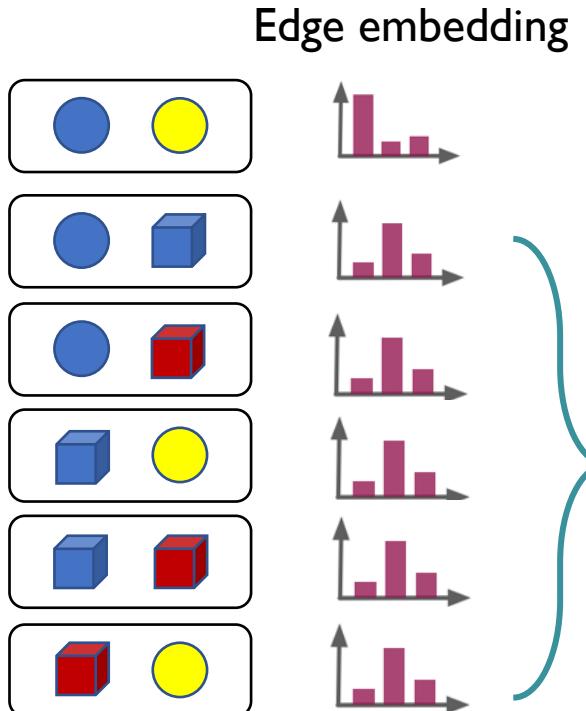
Kipf et al. Neural relational inference for interacting systems. ICML2018

## Relation inference for physical model

Observed trajectory



GNN encoder



Type I (Opposite charged)

Type 2 (No charge force)

- The charge property of objects is expressed by the charge relationship between objects
- Different relationship types reflect different types of forces
- Other types of forces such as springs and light rods can also be considered
- Relations are constant

# GNN for relation inference

Kipf et al. Neural relational inference for interacting systems. ICML2018

Accuracy (in %) of unsupervised interaction recovery

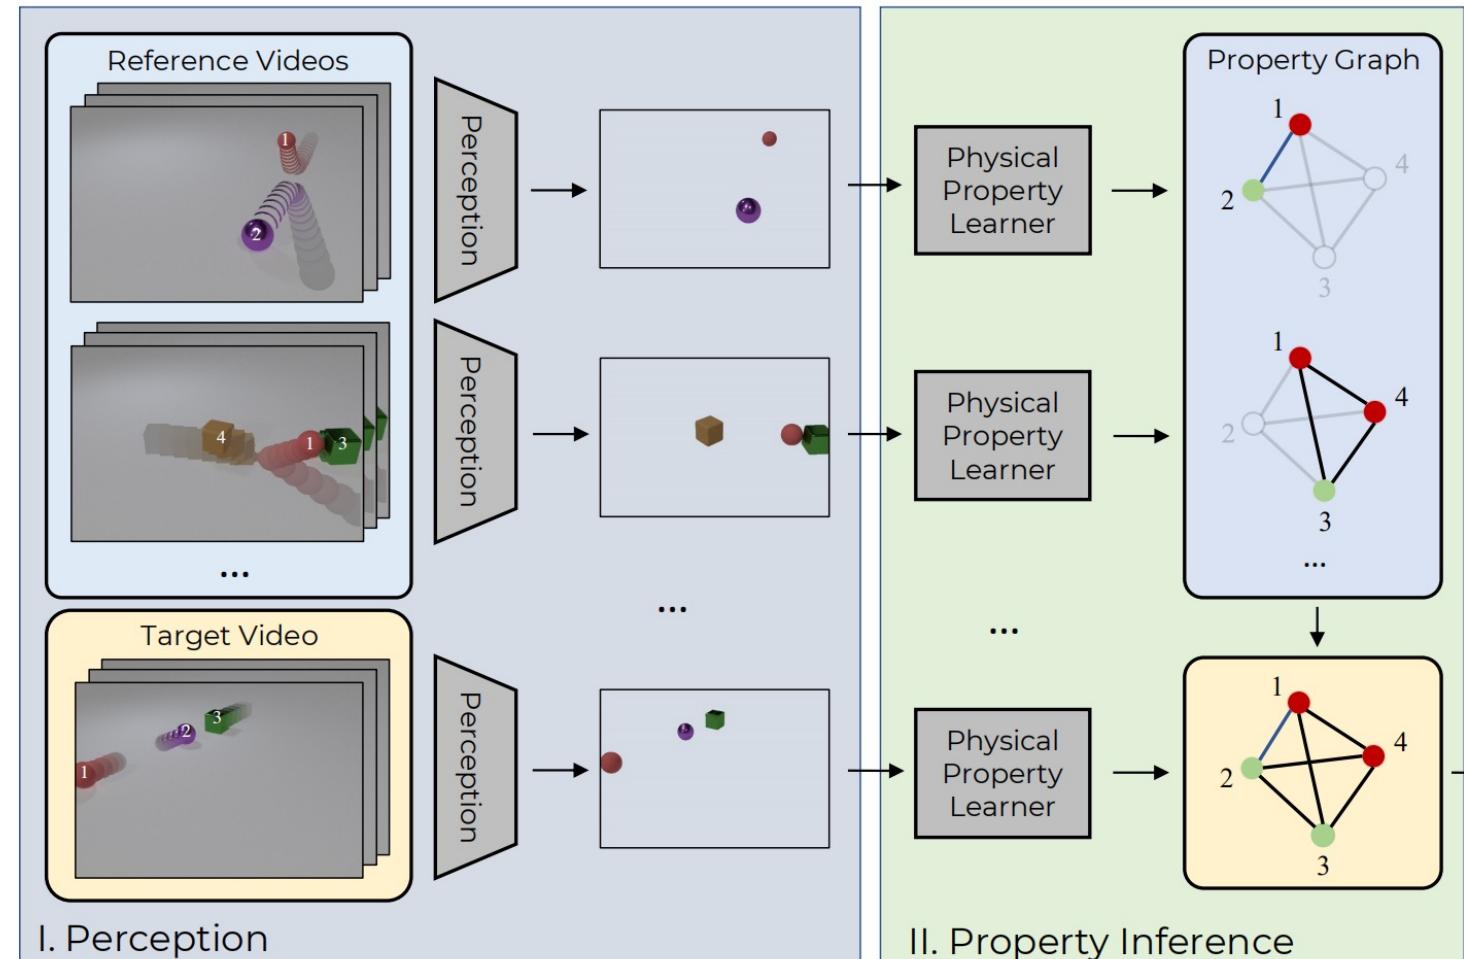
Model	Springs	Charged	Kuramoto
5 objects			
Corr. (path)	52.4 $\pm$ 0.0	55.8 $\pm$ 0.0	62.8 $\pm$ 0.0
Corr. (LSTM)	52.7 $\pm$ 0.9	54.2 $\pm$ 2.0	54.4 $\pm$ 0.5
NRI (sim.)	<b>99.8</b> $\pm$ 0.0	59.6 $\pm$ 0.8	—
NRI (learned)	<b>99.9</b> $\pm$ 0.0	<b>82.1</b> $\pm$ 0.6	<b>96.0</b> $\pm$ 0.1
Supervised	99.9 $\pm$ 0.0	95.0 $\pm$ 0.3	99.7 $\pm$ 0.0
10 objects			
Corr. (path)	50.4 $\pm$ 0.0	51.4 $\pm$ 0.0	59.3 $\pm$ 0.0
Corr. (LSTM)	54.9 $\pm$ 1.0	52.7 $\pm$ 0.2	56.2 $\pm$ 0.7
NRI (sim.)	<b>98.2</b> $\pm$ 0.0	53.7 $\pm$ 0.8	—
NRI (learned)	<b>98.4</b> $\pm$ 0.0	<b>70.8</b> $\pm$ 0.4	<b>75.7</b> $\pm$ 0.3
Supervised	98.8 $\pm$ 0.0	94.6 $\pm$ 0.2	97.1 $\pm$ 0.1

- Recovers the ground-truth interaction graph with high accuracy on most tasks.
- Unsupervised method of NRI is comparable to the supervised benchmark

# GNN for relation inference

Chen et al. COMPHY: Compositional physical reasoning of objects and events from videos. ICLR 2022

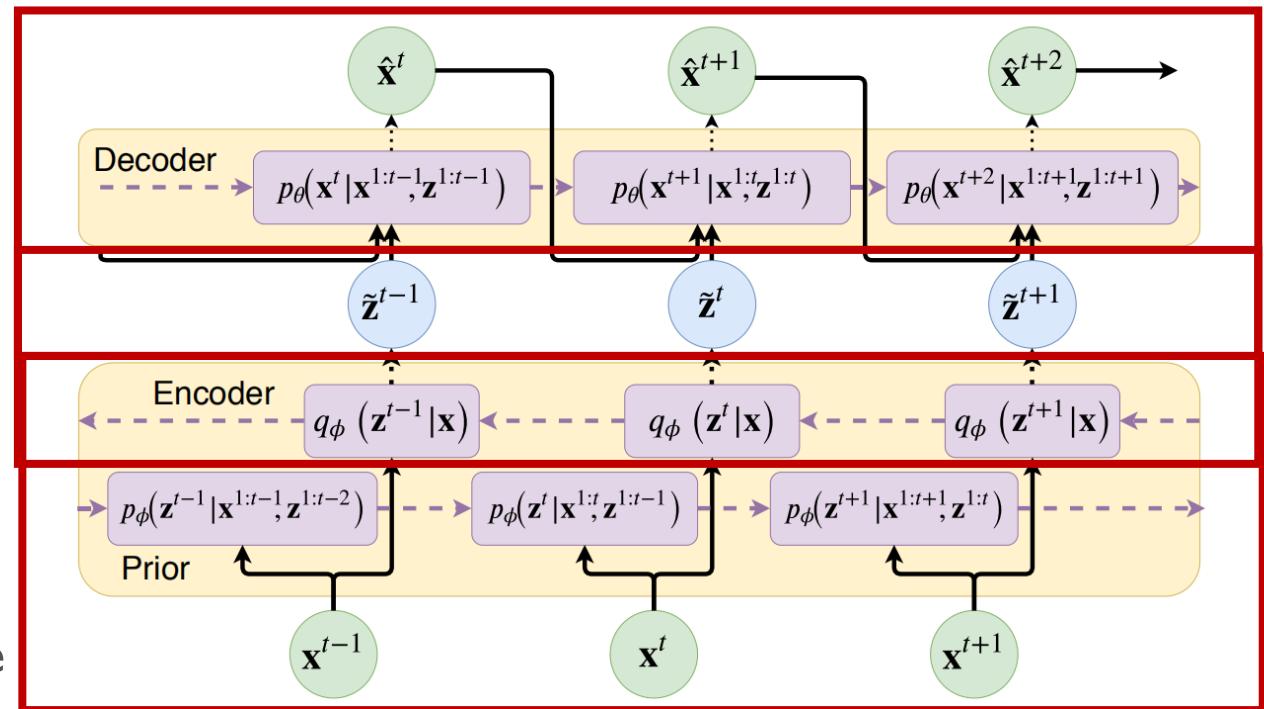
- Latest dataset **COMPHY** provides 4 reference videos for each target video.
- The object in the reference video has appeared in the target video.
- Property graph is expected to be better inferred with reference videos.



# GNN for Dynamic Neural Relational Inference

Graber, Schwing. Dynamic Neural Relational Inference. CVPR 2020

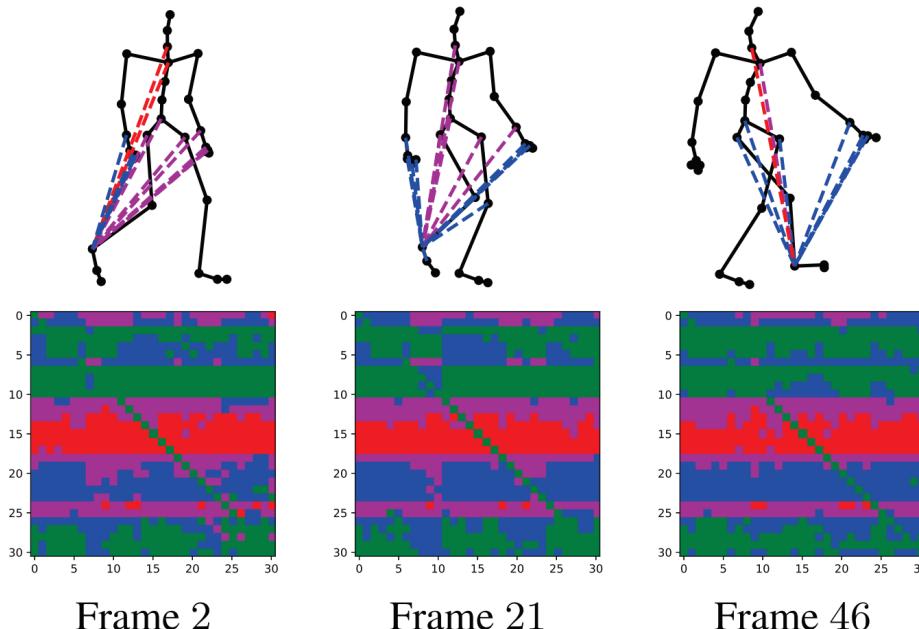
- Produce edge embeddings at every time step
- LSTM encodes the past and future history of entity relations
- Approximate posterior is computed as a function of both the past and future.
- Edge variables are sampled from the approximate posterior  
(to select edge models for the decoder)
- The decoder evolves a hidden state using this GNN and the previous predictions
- Predicts the state of the entities at the next time step.



# GNN for Dynamic Neural Relational Inference

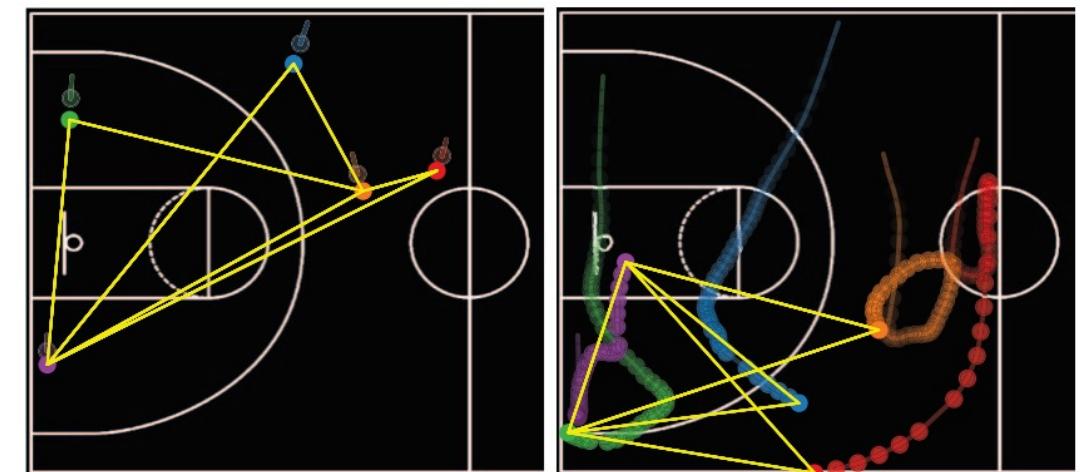
Graber, Schwing. Dynamic Neural Relational Inference. CVPR 2020

Motion prediction



Different motion stages imply different motion relationships of joint points

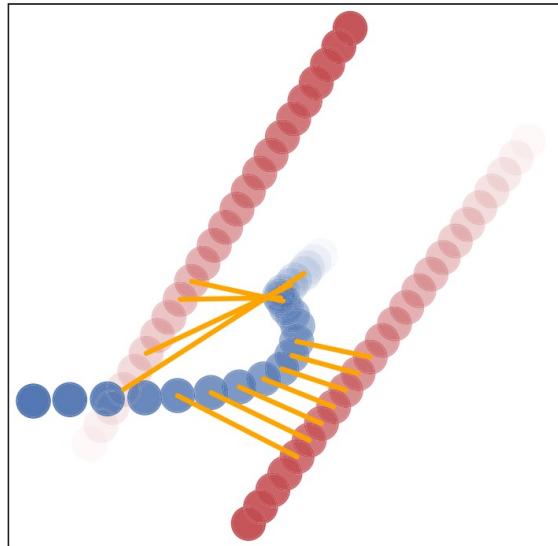
NBA basketball game dataset



The players' coordination changes over time

# GNN for Dynamic Neural Relational Inference

Graber, Schwing. Dynamic Neural Relational Inference. CVPR 2020

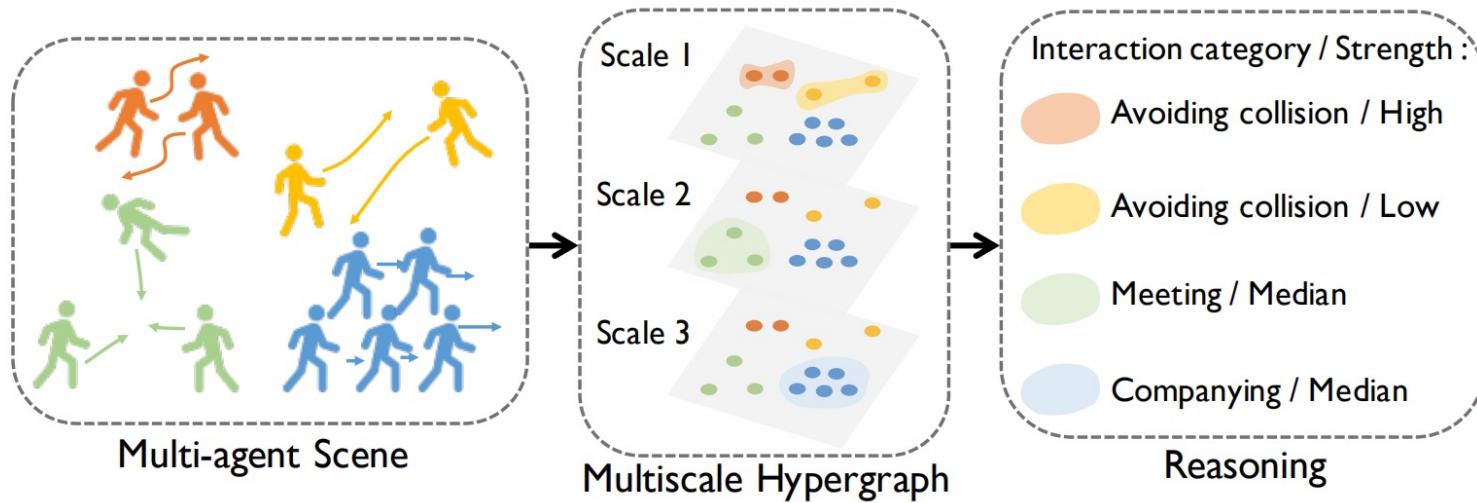


- The first two (red) move with a constant velocity in some direction.
- The third (blue) is initialized with a random velocity but is additionally “pushed” away by the other particles whenever the distance separating them is less than 1. (Such high-frequency function is not easy to fit in static NRI)
- Dynamic relationships reflect dynamic forces

Dynamic relation in physical dataset

# GNN for hyper-relational Inference

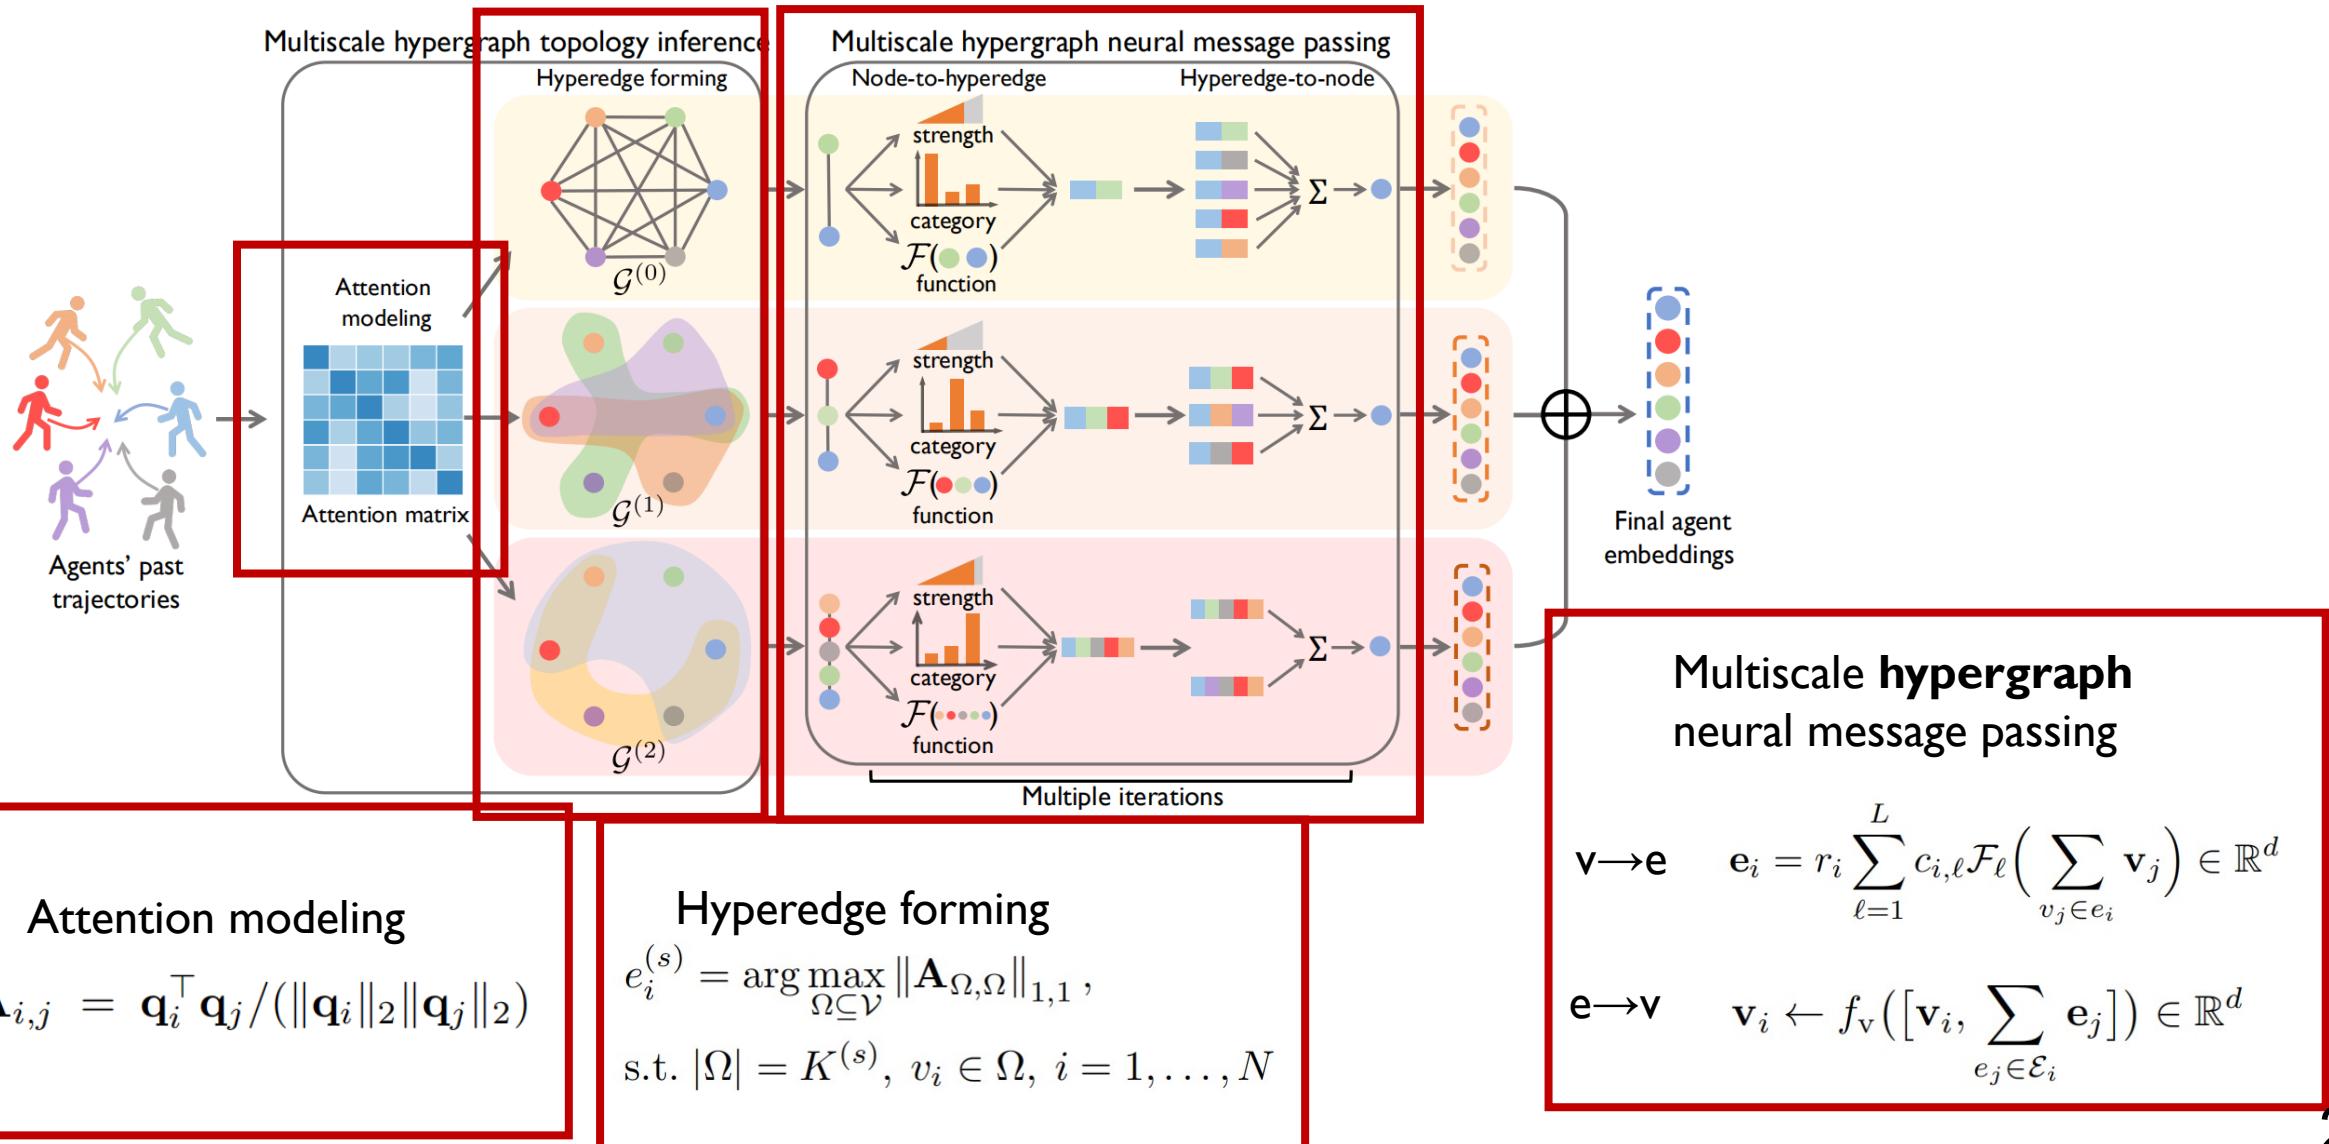
Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022



- Captures both pair-wise and group-wise interactions among multiple agents from their trajectories.
- Infers the interaction category and strength for each interaction group.

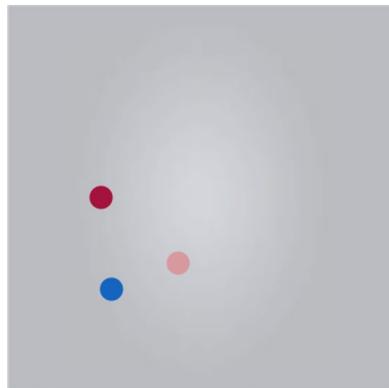
# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022

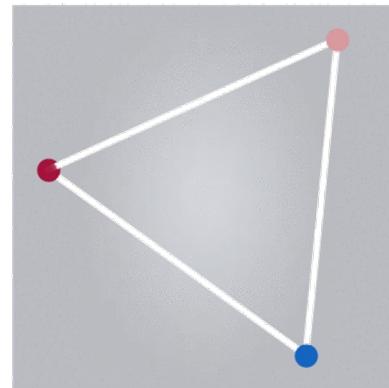


# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022



Free



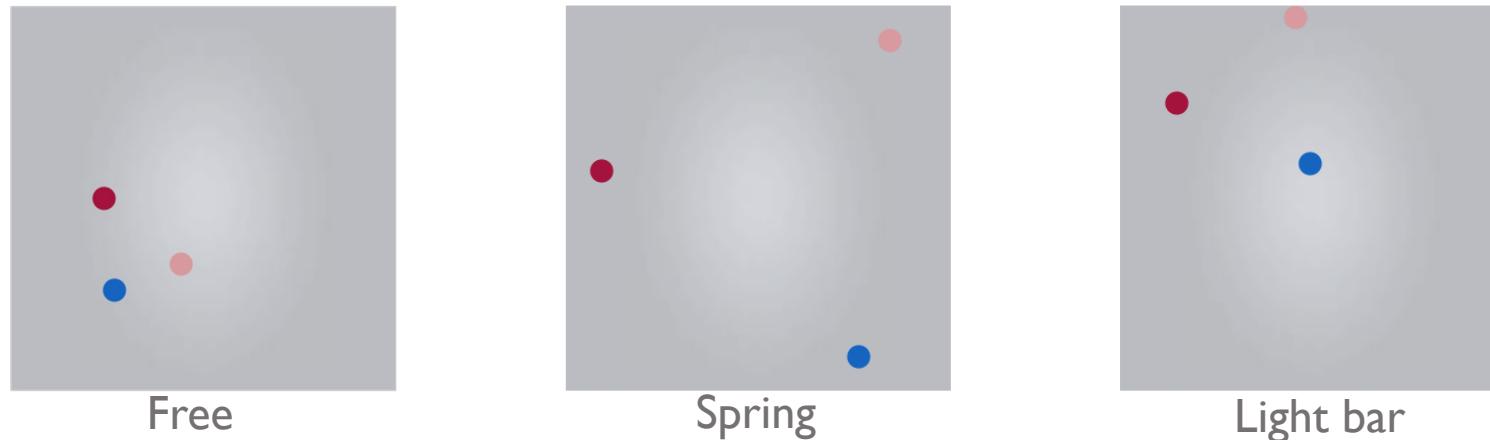
Spring



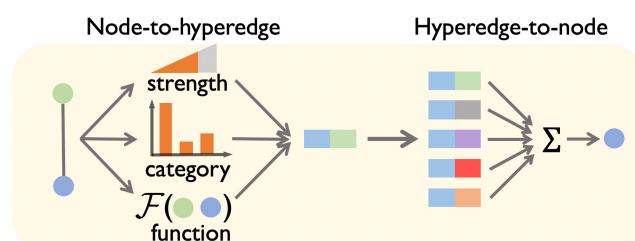
Light bar

# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022



**Task:** Given the trajectories **only**, recognize the interaction categories during prediction

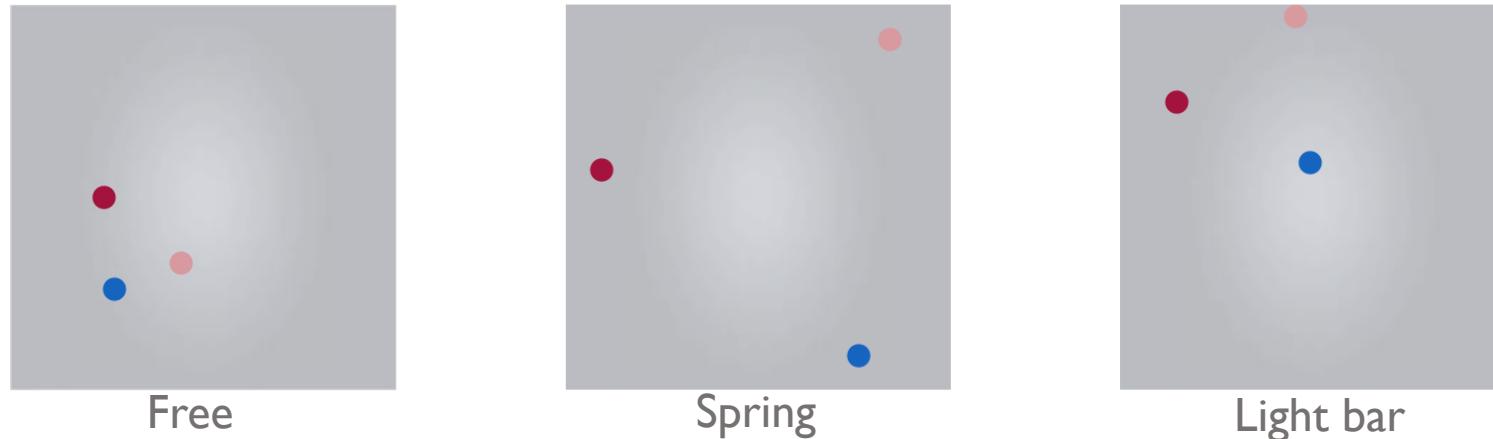


$$\mathbf{e}_i = r_i \sum_{\ell=1}^L c_{i,\ell} \mathcal{F}_\ell \left( \sum_{v_j \in e_i} \mathbf{v}_j \right) \in \mathbb{R}^d,$$

↑  
category

# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022



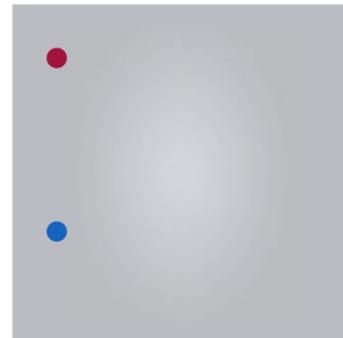
**Task:** Given the trajectories **only**, recognize the interaction categories during prediction

Table 1: Comparison of accuracy (mean  $\pm$  std in %) of on interaction category recognition.

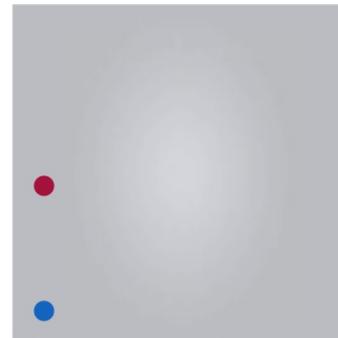
	Corr. (path)	Corr. (LSTM)	NRI	Ours	Supervised
2-types	78.3 $\pm$ 0.0	75.3 $\pm$ 2.5	89.7 $\pm$ 3.5	<b>99.3 <math>\pm</math> 0.3</b>	99.5 $\pm$ 0.2
3-types	-	-	64.9 $\pm$ 2.3	<b>91.5 <math>\pm</math> 2.4</b>	98.7 $\pm$ 0.7

# GNN for hyper-relational Inference

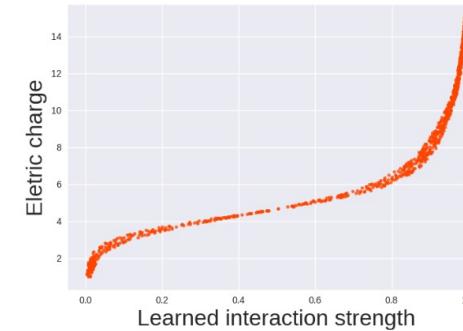
Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022



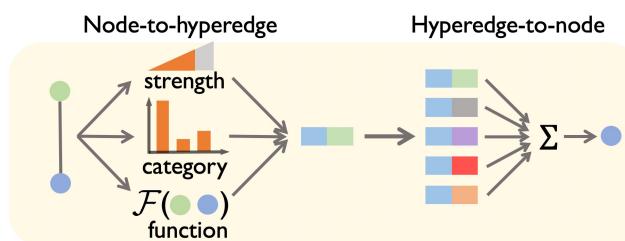
Low charge



High charge



Strength-Charge Curve



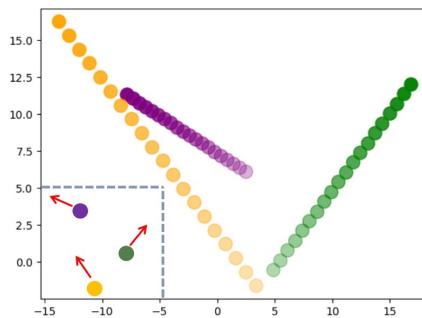
$$\mathbf{e}_i = r_i \sum_{\ell=1}^L c_{i,\ell} \mathcal{F}_\ell \left( \sum_{v_j \in e_i} \mathbf{v}_j \right) \in \mathbb{R}^d,$$

↑  
strength

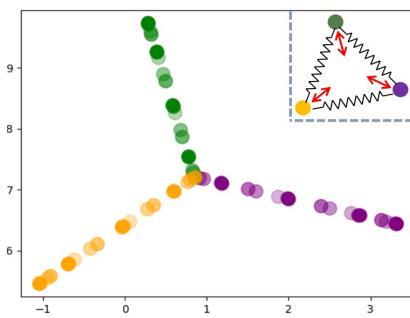
# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022

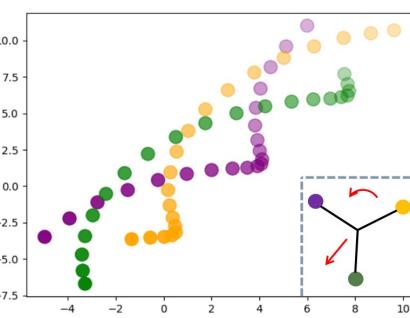
Trajectory samples under three interaction categories.



(a) Free

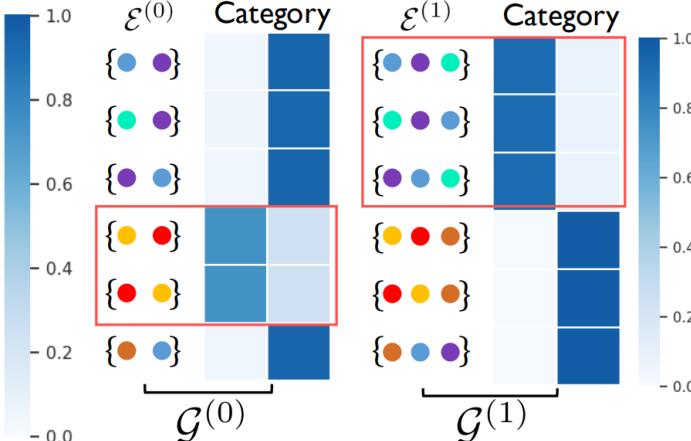
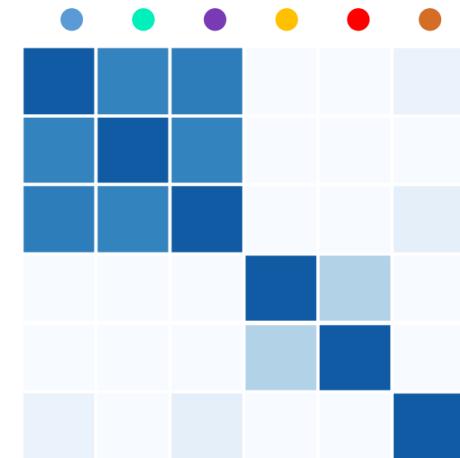
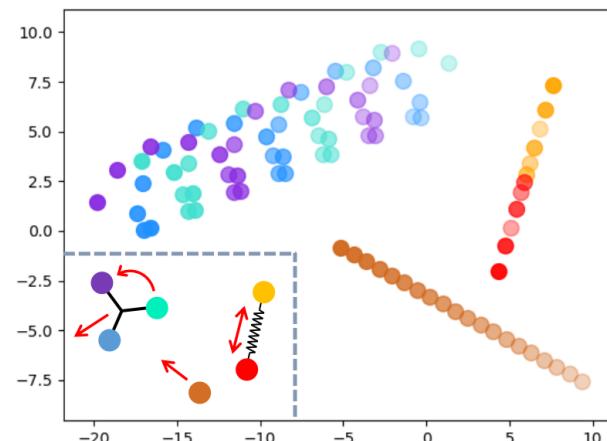


(b) Spring



(c) Light bar

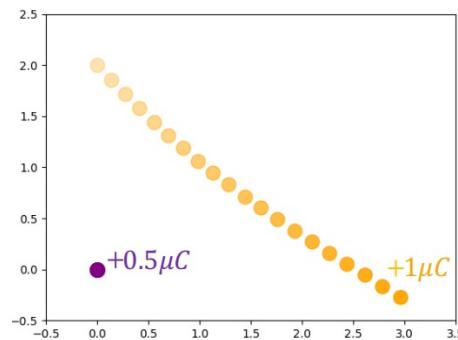
- The particle trajectories contain a three-group with a light bar, a two-group with a spring and an individual particle.
- The heatmap of the learnt affinity matrix via affinity modeling.
- The red box represents the three-group and two-group we inferred.



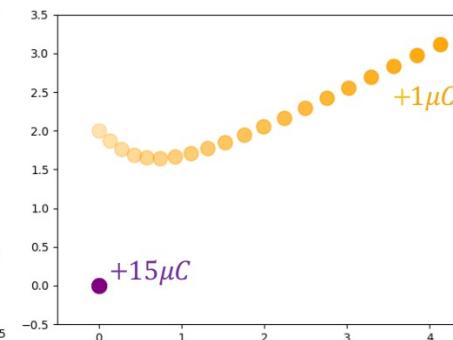
# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022

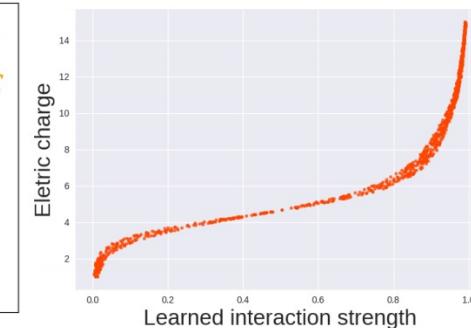
- Both category and strength of interaction can be captured



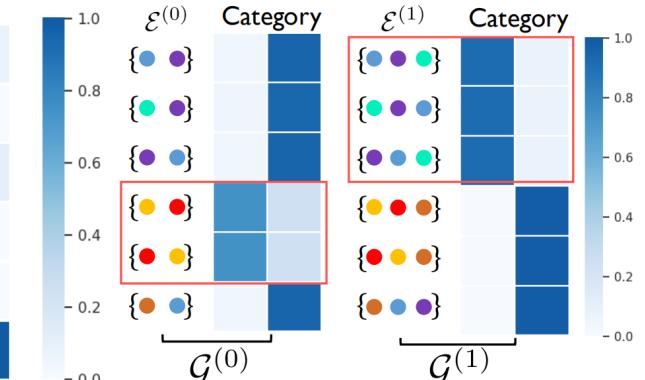
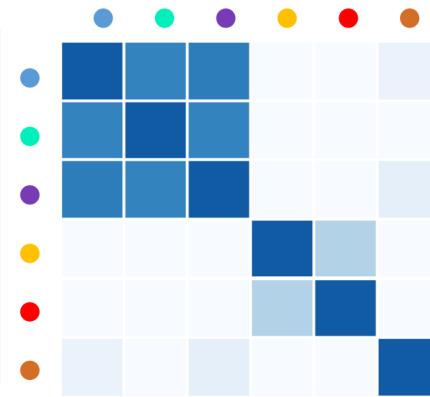
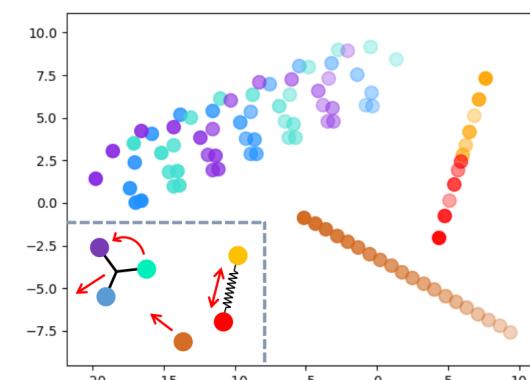
(a) Low charge



(b) High charge



(c) Relation curve

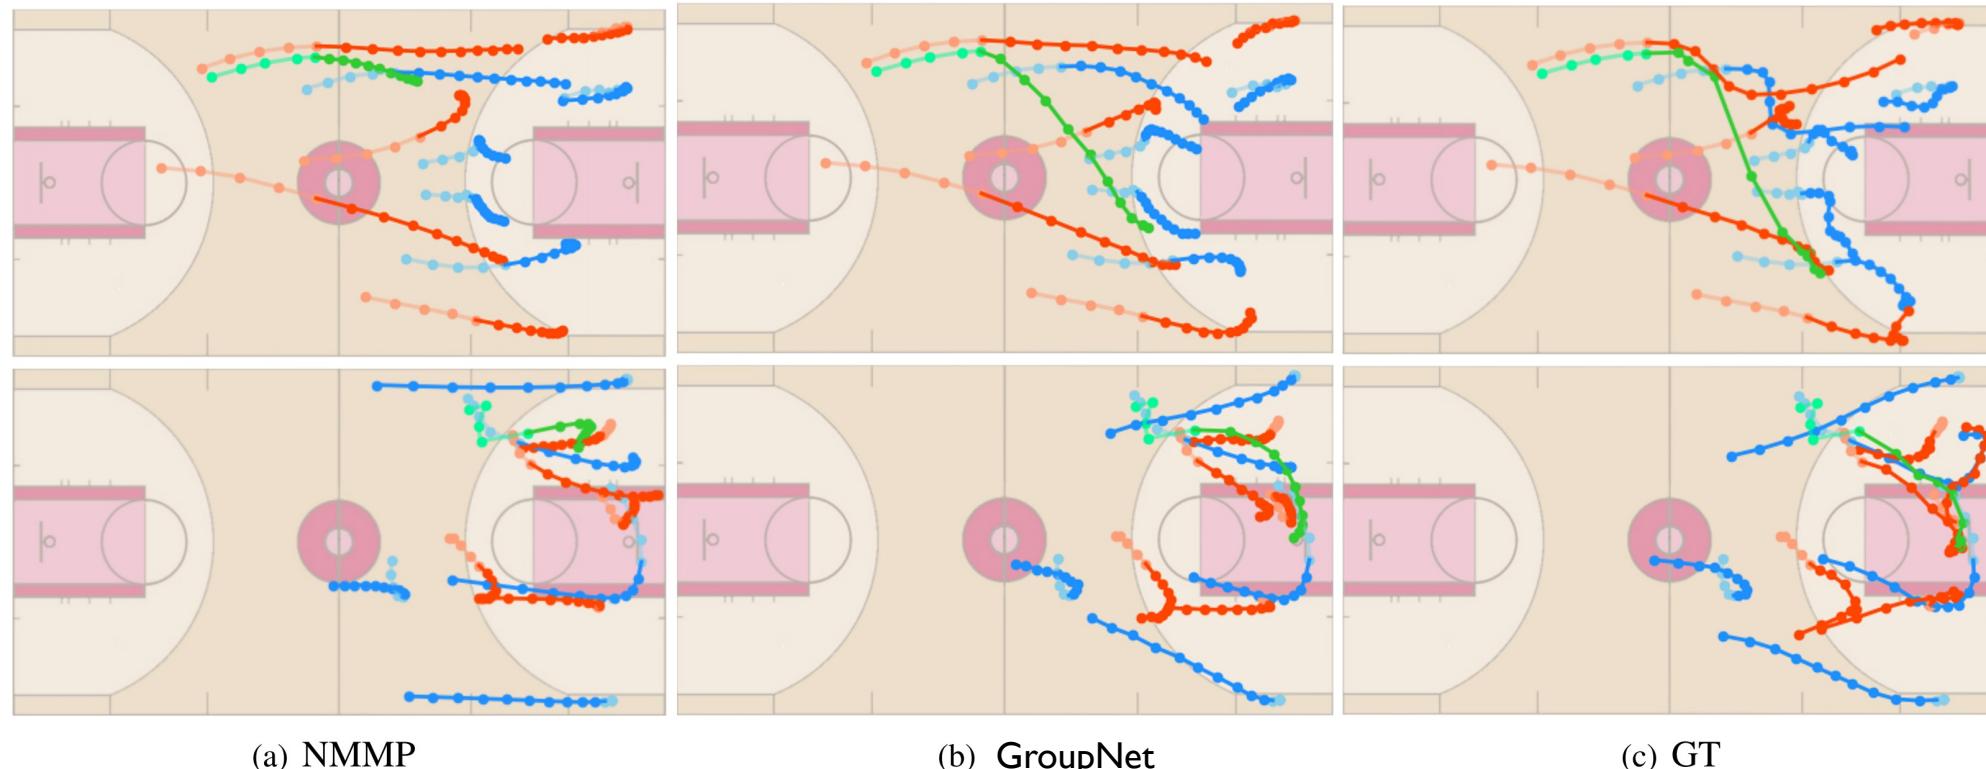


# GNN for hyper-relational Inference

Xu, Li, Ni, Zhang, Chen. GroupNet. CVPR 2022

Qualitative results on the NBA dataset

- The best trajectory among 20 predictions
- Red/blue — two teams
- Green — basketball
- Light color — past trajectory



# GNN for complex Physics

Alvaro et al. Learning to Simulate Complex Physics with Graph Networks. ICML 2020

GNN for a wide variety of challenging

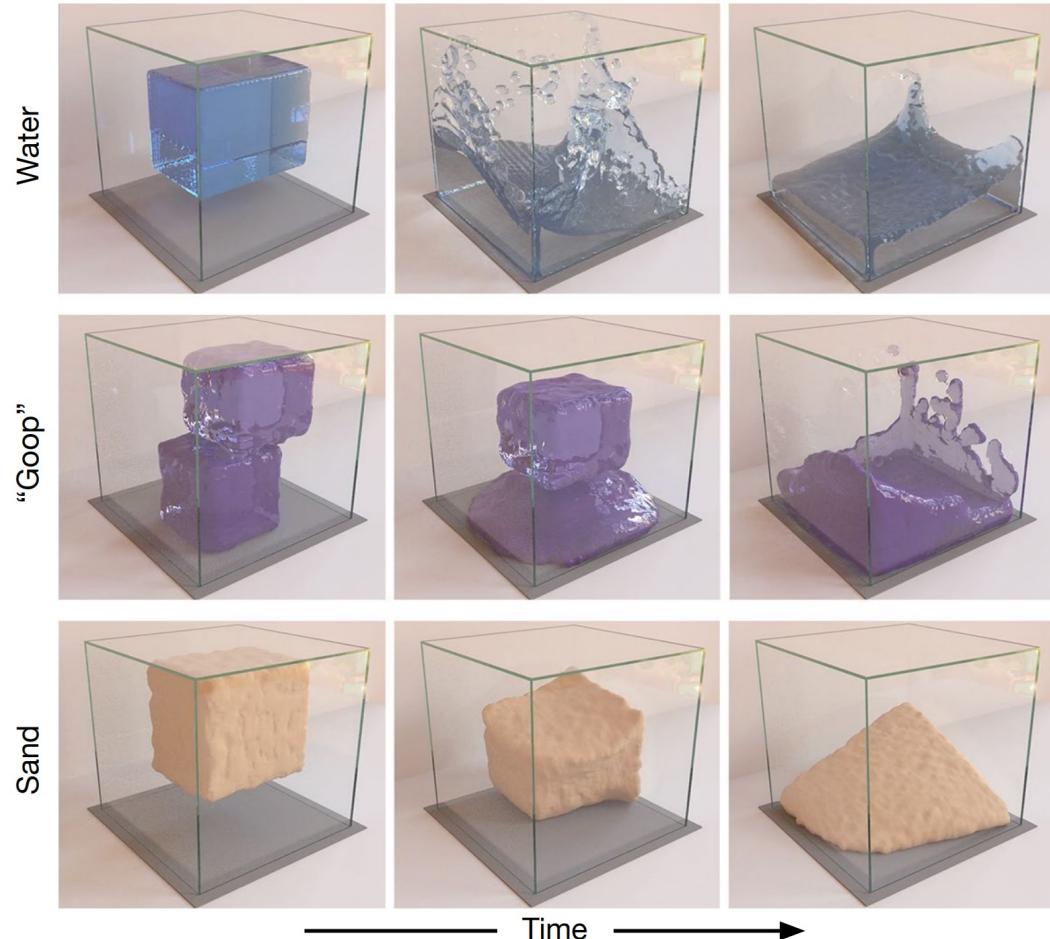
physical domains:

- Fluids
- Rigid solids
- deformable materials

Physical model—thousands of particles

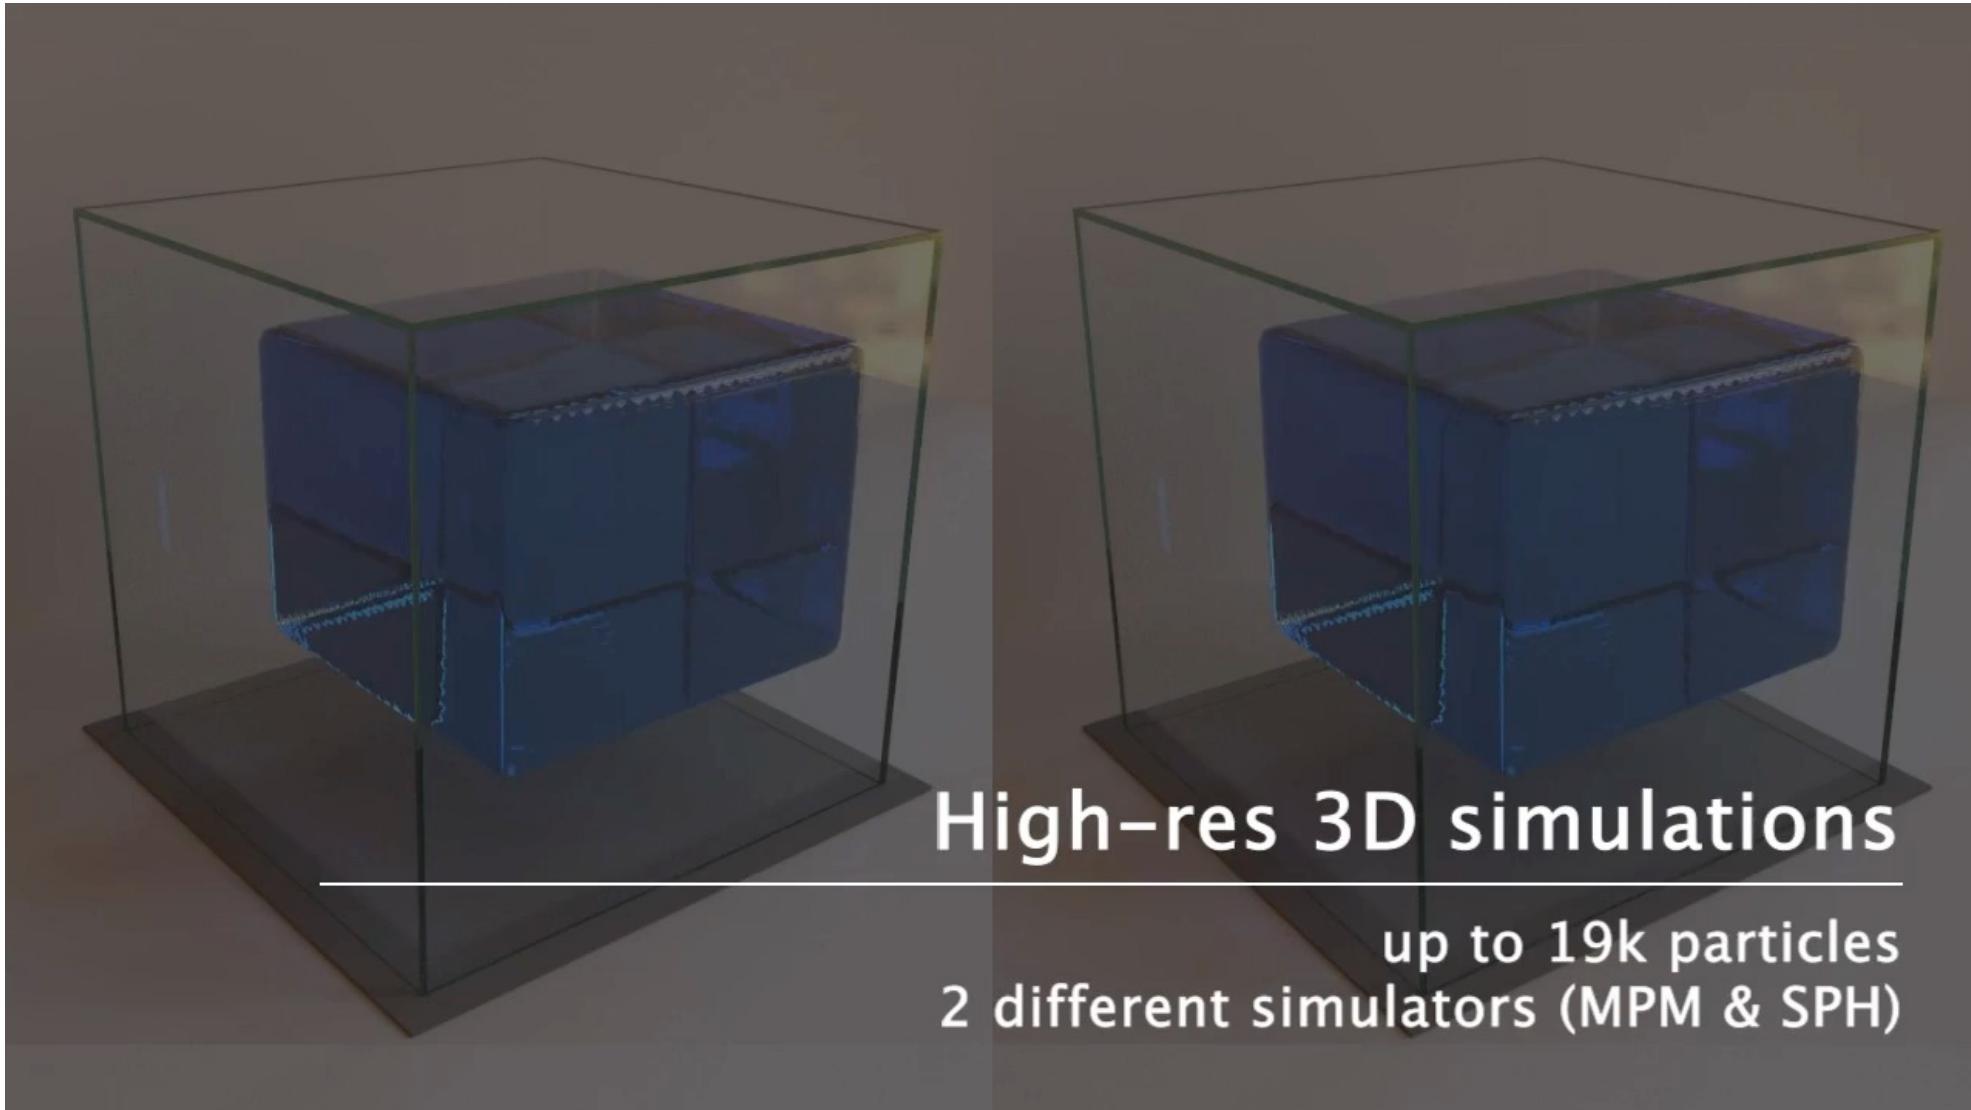
Particles—nodes in a graph,

Dynamics—message-passing for prediction



# GNN for complex Physics

Alvaro et al. Learning to Simulate Complex Physics with Graph Networks. ICML 2020

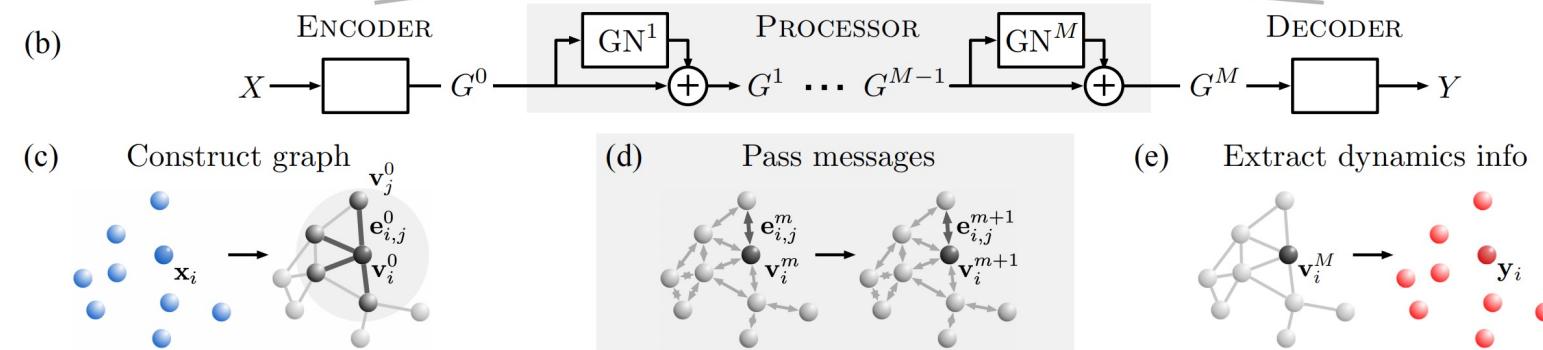
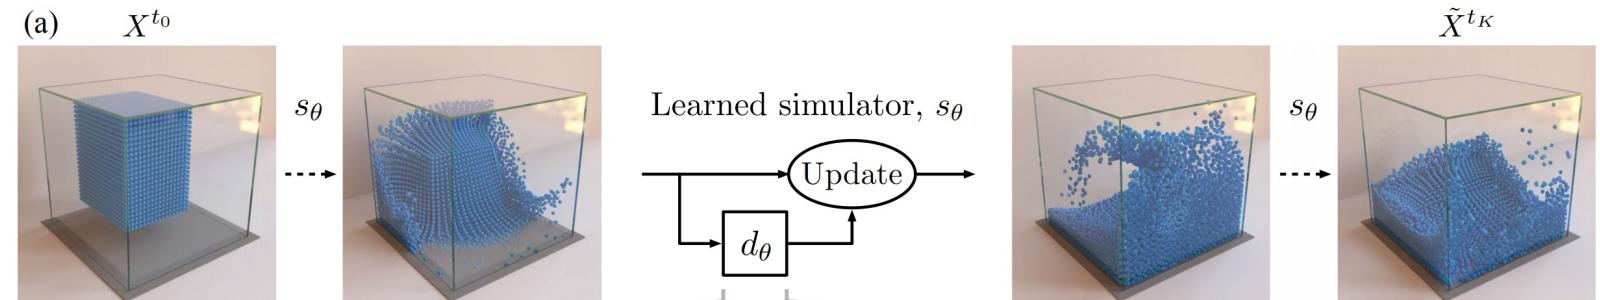


# GNN for complex Physics

Alvaro et al. Learning to Simulate Complex Physics with Graph Networks. ICML 2020

## Encode-process-decode scheme:

- Encoder: constructs latent graph  $G$  from input state  $X$
- Processor: performs  $M$  rounds of learned message-passing over the latent graphs  $G^1, G^2, \dots, G^M$
- Decoder: extracts dynamics information from the final latent graph  $G^M$



Add edges between particles within a “connectivity radius” (local interactions prior)

Multi-step message-passing is critical for complex physics (especially for rigid object)

Future position and velocity are updated

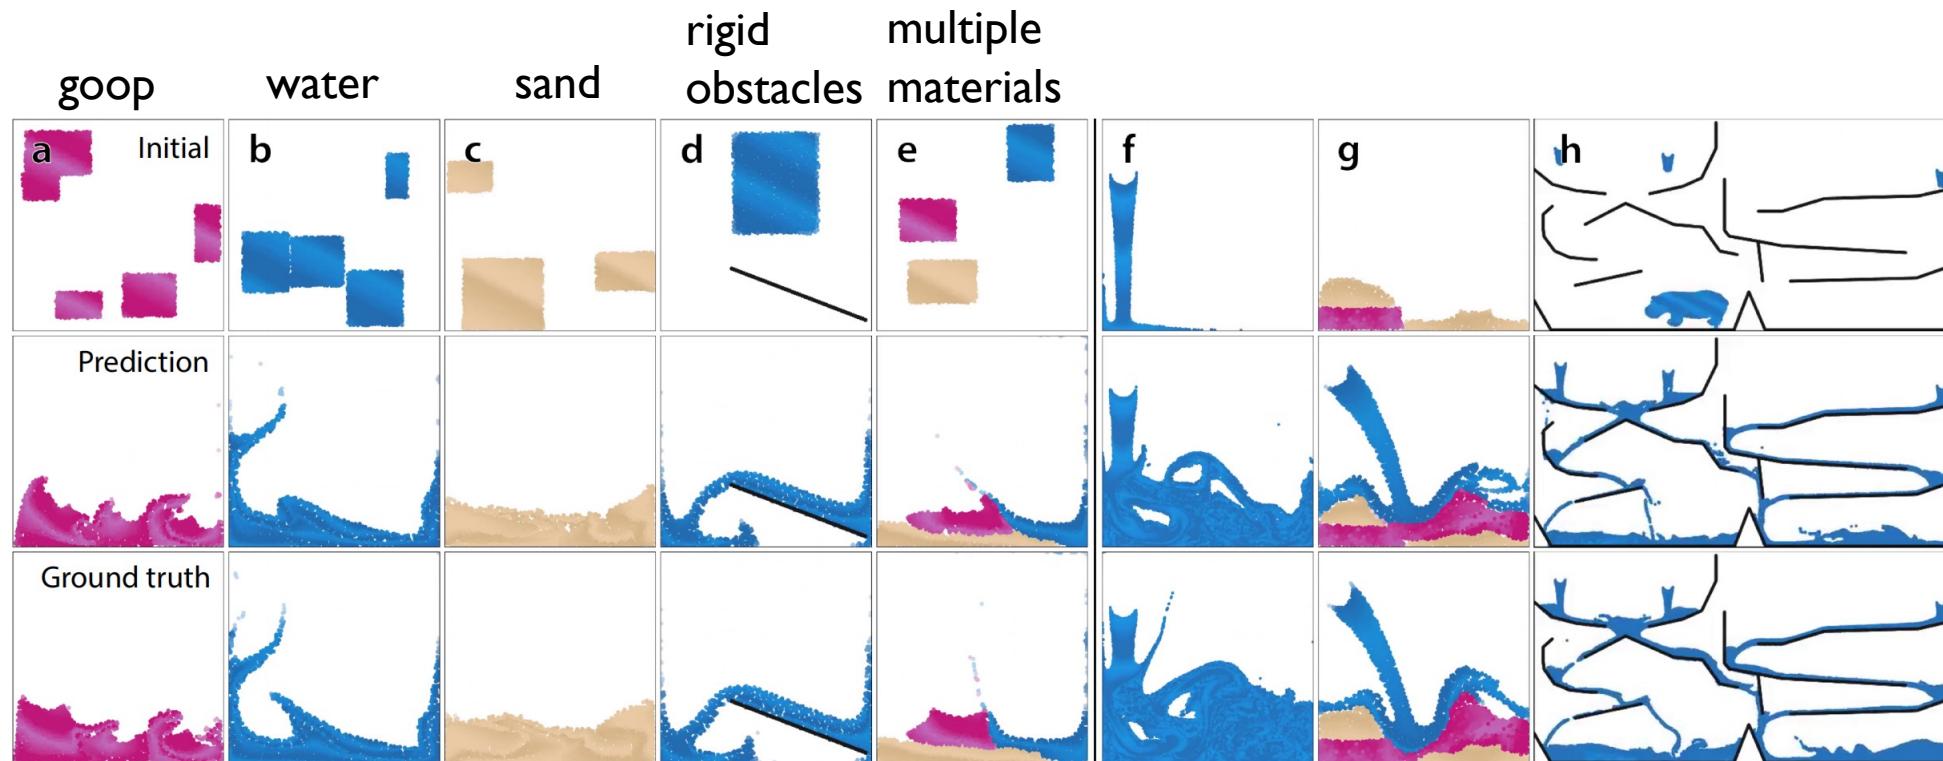
# GNN for complex Physics

Alvaro et al. Learning to Simulate Complex Physics with Graph Networks. ICML 2020

## Simulate many materials

Pre-trained models on several out-of-distribution:

- Water ramps → high-res turbulence (f)
- Multi-material interactions → Multi-material interactions +unseen objects (g)
- generalizing on significantly larger domains (h)



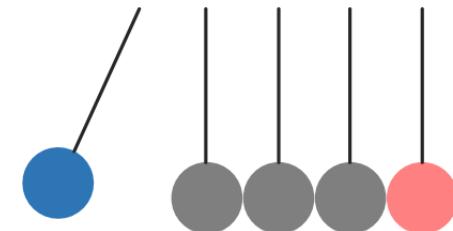
# Hierarchical modeling for dynamic particle interaction

Li et al. Learning particle dynamics for manipulating rigid bodies, deformable objects, and fluids. ICLR 2019

## Modeling of the long-range dependence via hierarchical modeling

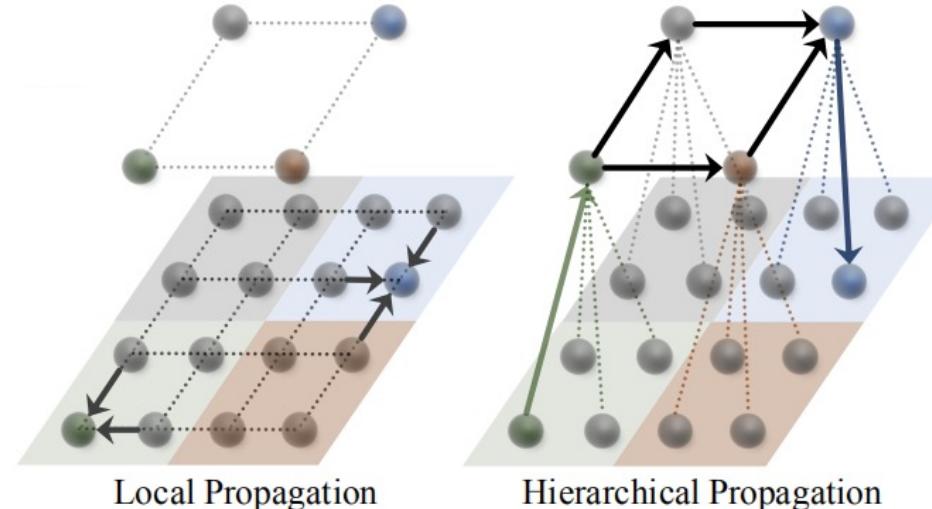
- A object consists of particles (leaf).
- Cluster the particles into several non-overlapping clusters.
- A new root particle for each cluster.

For handling instantaneous force propagation (long-range dependence)

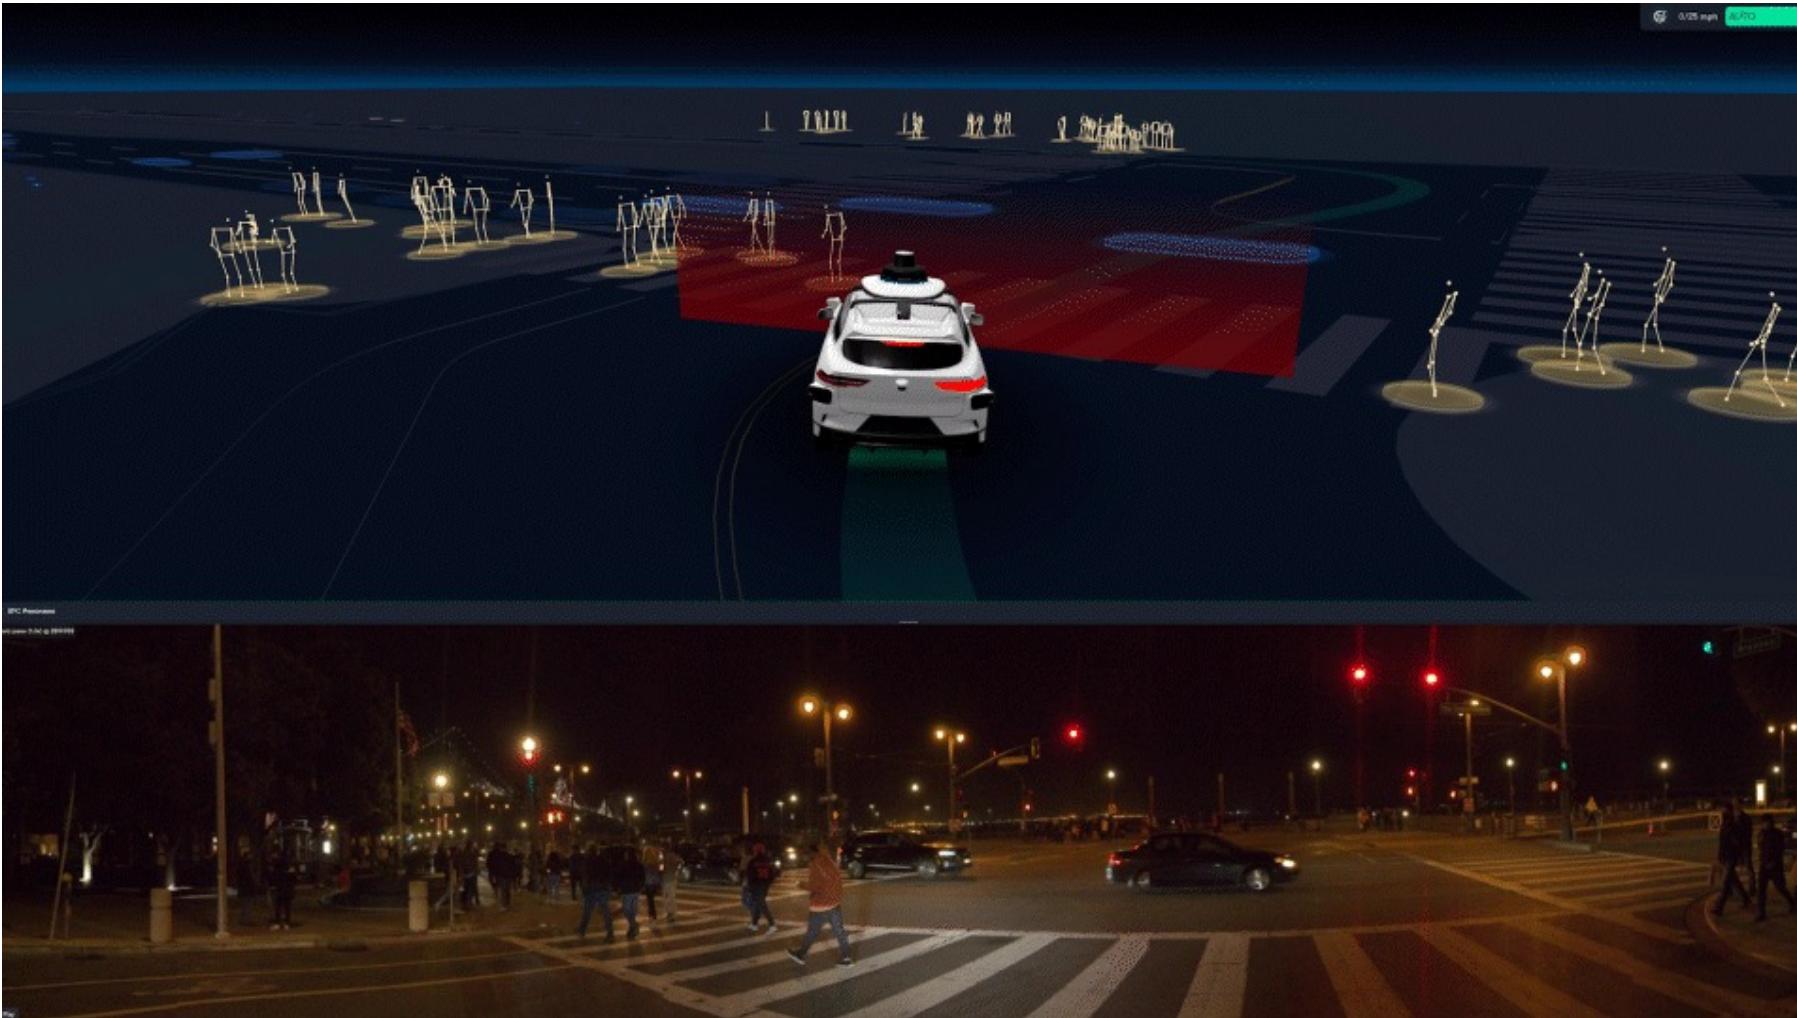


## Multi-stage propagation

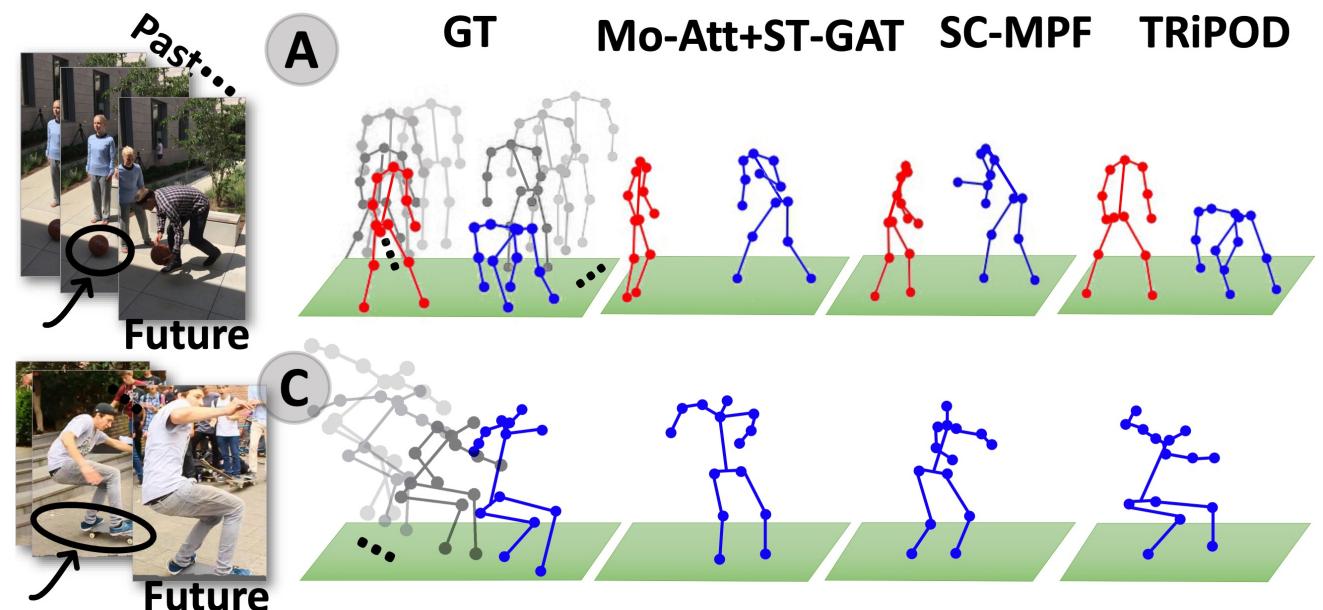
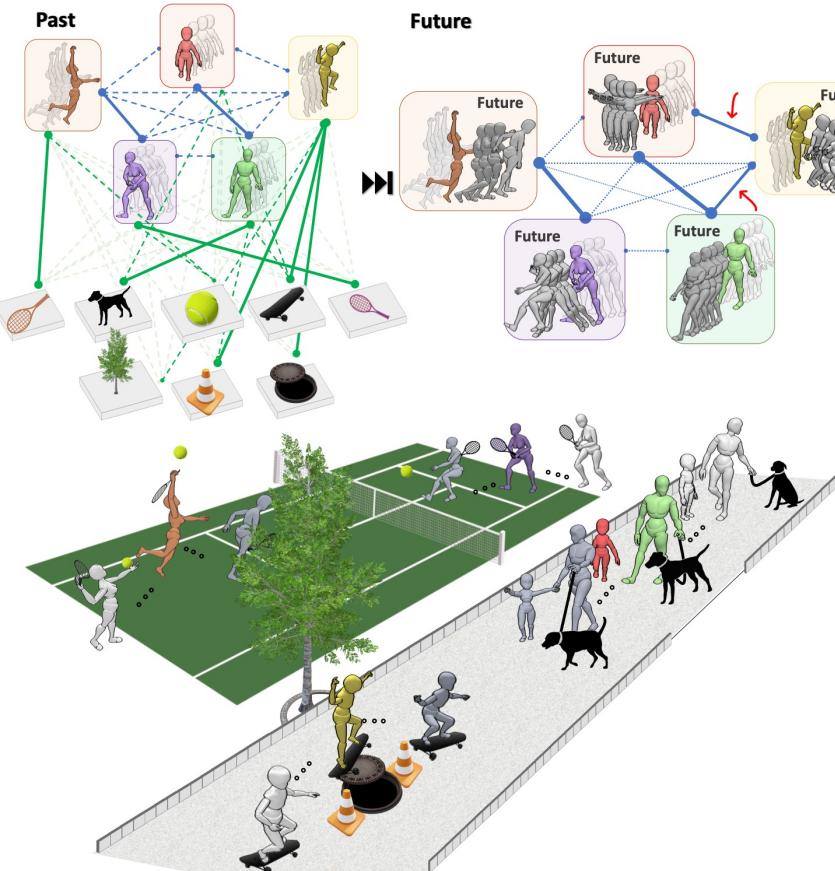
- Leaf → Leaf
- Leaf → Root
- Root → Root
- Root → Leaf



# GCN for Multi-Person Behavior Prediction

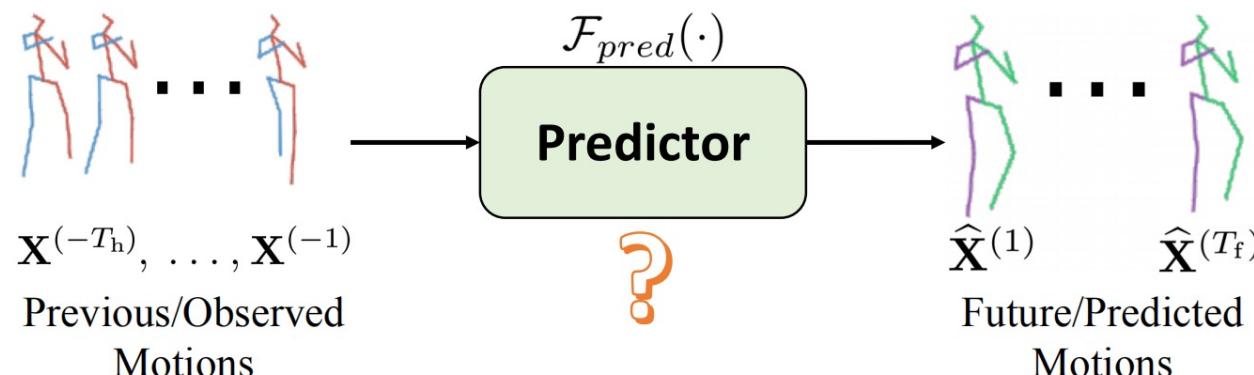


# GCN for Multi-Person Behavior Prediction

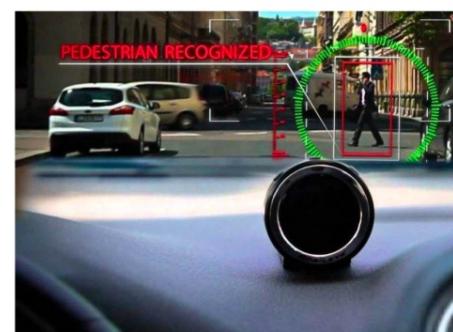


# GCN for Skeleton based Human Motion Prediction

**What is 3D skeleton-based human motion prediction?**



Human-machine Interaction



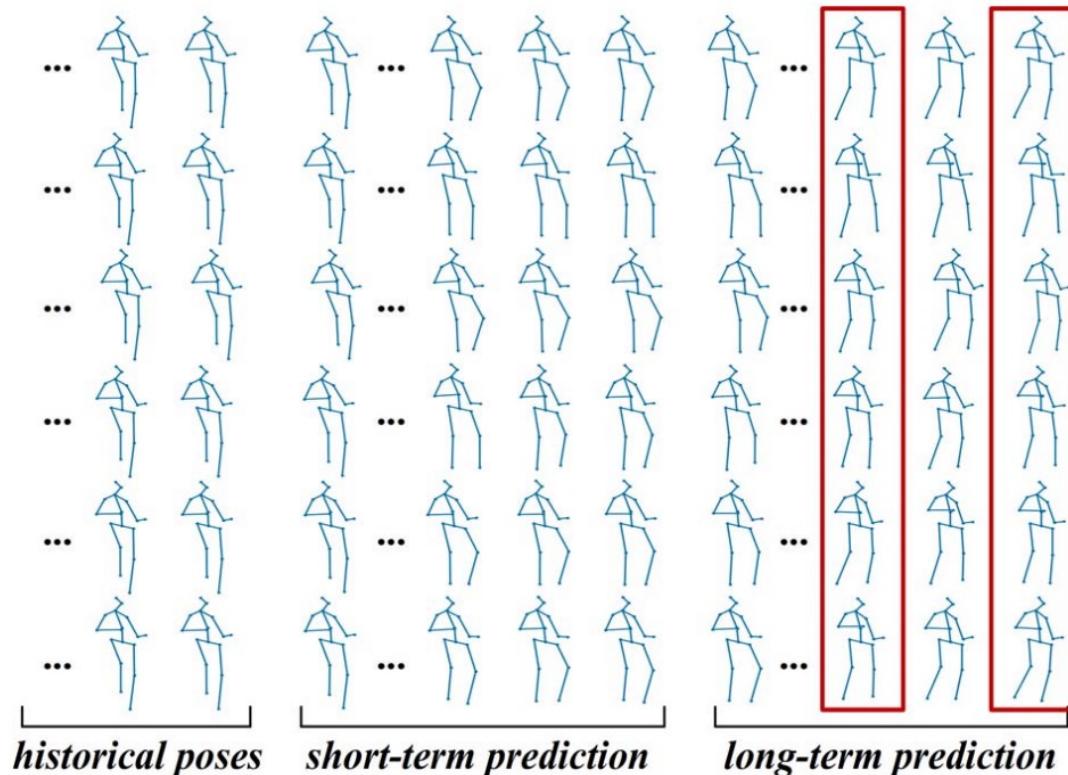
Automatic driving



Tracking and Surveillance

# GCN for Skeleton based Human Motion Prediction

## Short-term/long-term motion prediction



Varies of motion type  
(H3.6M dataset):

- Walking
- Eating
- Smoking
- Discussion
- Directions
- Greeting
- Phoning
- Posing
- Purchase
- Sitting
- ...

Temporal modeling:

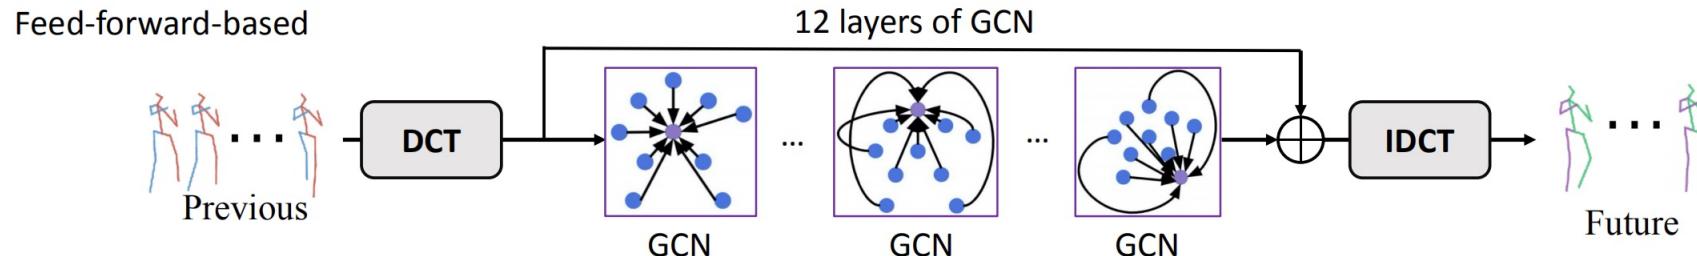
- Time convolution
- DCT coding

Skeleton modeling:

- Graph convolution network

# GCN for Skeleton based Human Motion Prediction

Cui et al. Learning Dynamic Relationships for 3D Human Motion Prediction. CVPR 2020



Temporal modeling:

- Time convolution
- DCT coding

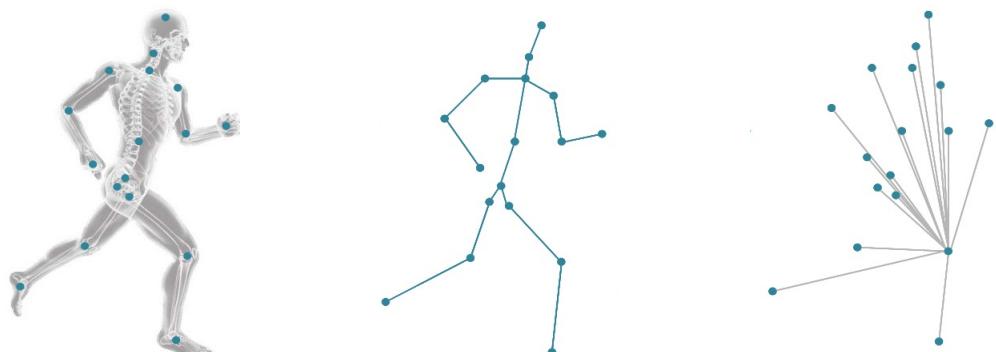
Skeleton modeling:

- Graph convolution network

**Graph convolution:**  $\mathbf{F}^{(l+1)} = g(\mathbf{F}^{(l)}) = \sigma[(\mathbf{A}_p^{(l)} \circ \mathbf{M} + \mathbf{Q}^{(l)})\mathbf{F}^{(l)}\mathbf{W}^{(l)}],$

partially constrained

flexible



$\mathbf{F}^{(l)}$ : feature in layer  $l$

$\mathbf{M}$ : skeleton mask matrix

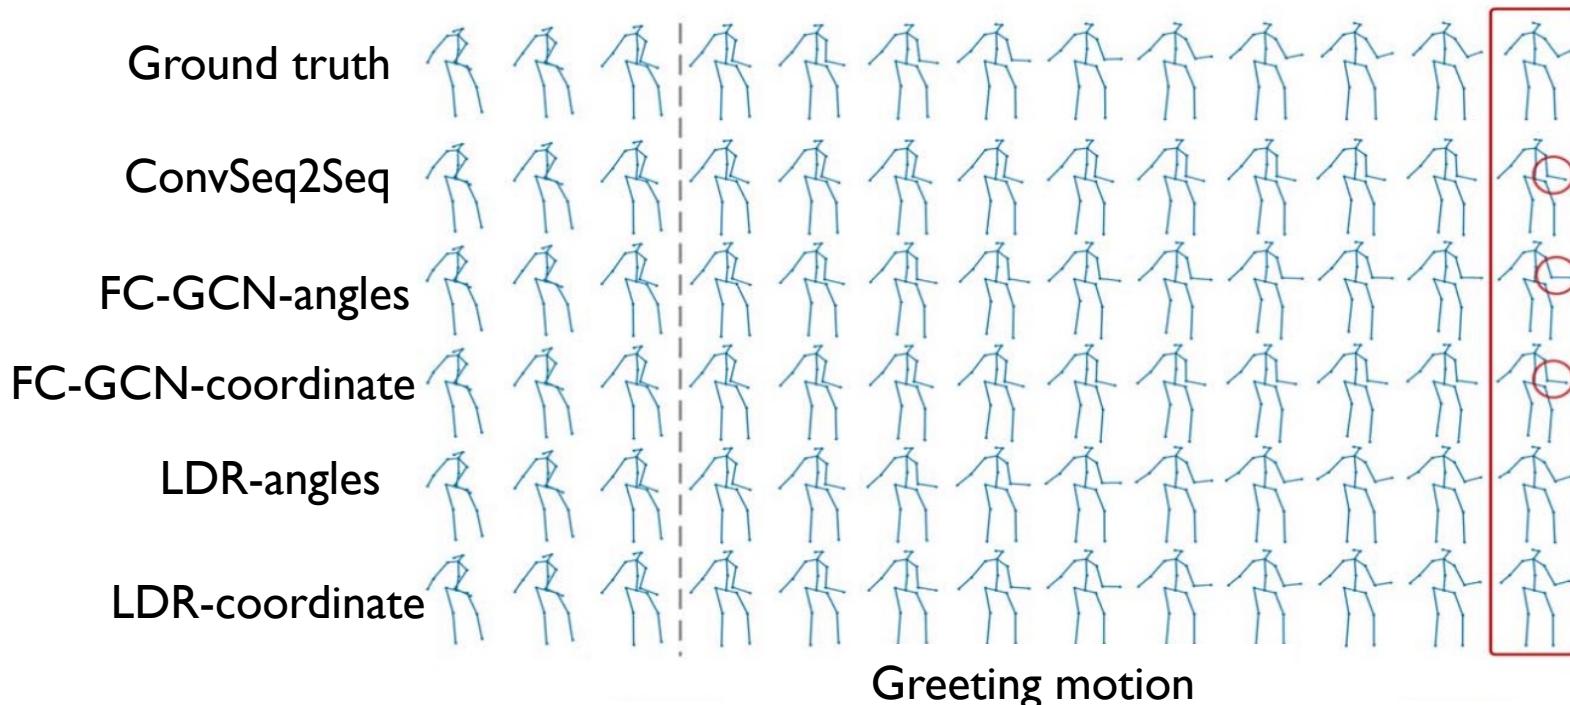
$\mathbf{W}$ : learnable weight matrix

## Over-smoothing

- Deep low-pass filtering
- Possible consequence:  
The speed of each joint point tends to be consistent, and the posture does not change

# GCN for Human Motion Prediction

Cui et al. Learning Dynamic Relationships for 3D Human Motion Prediction. CVPR 2020



Motion position coding:

- 3D coordinate (x,y,z)
- Angle of each joint

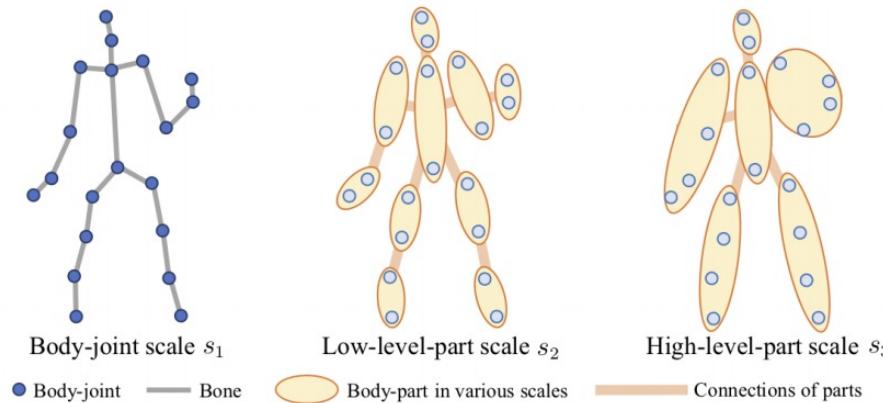
The red box represents the prediction result of the last frame (400 ms)

The red circle represents the wrong part of the model

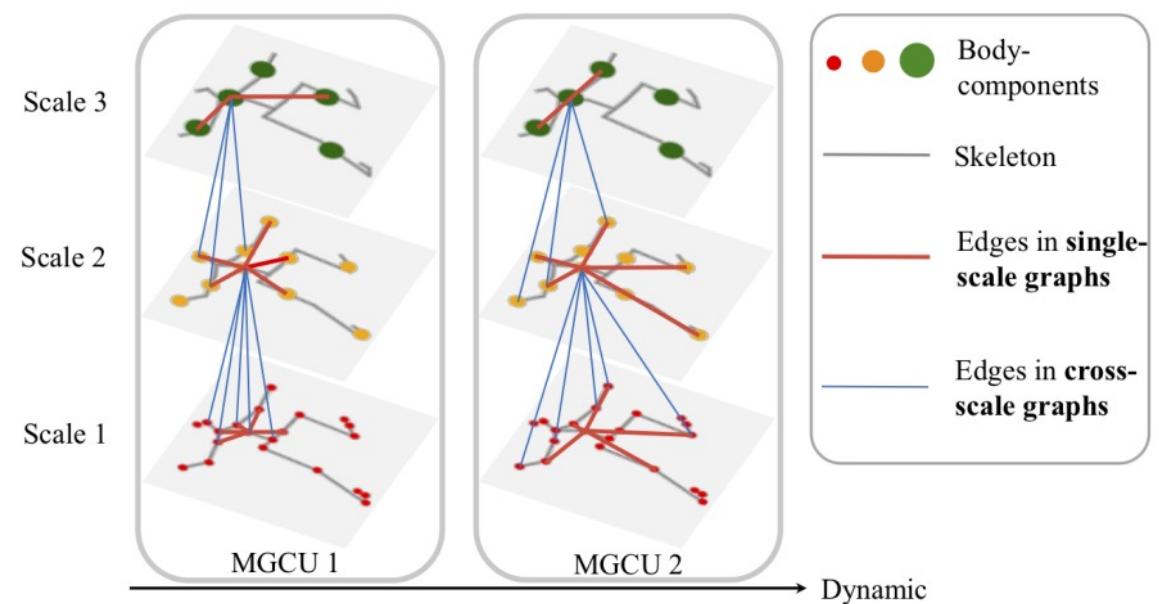
# Multiscale GCN model for Human Motion Prediction

Li, Chen, et al. Dynamic Multiscale Graph Neural Networks for 3D skeleton-based Human Motion Prediction. CVPR 2020 (oral)

**Three body scales**



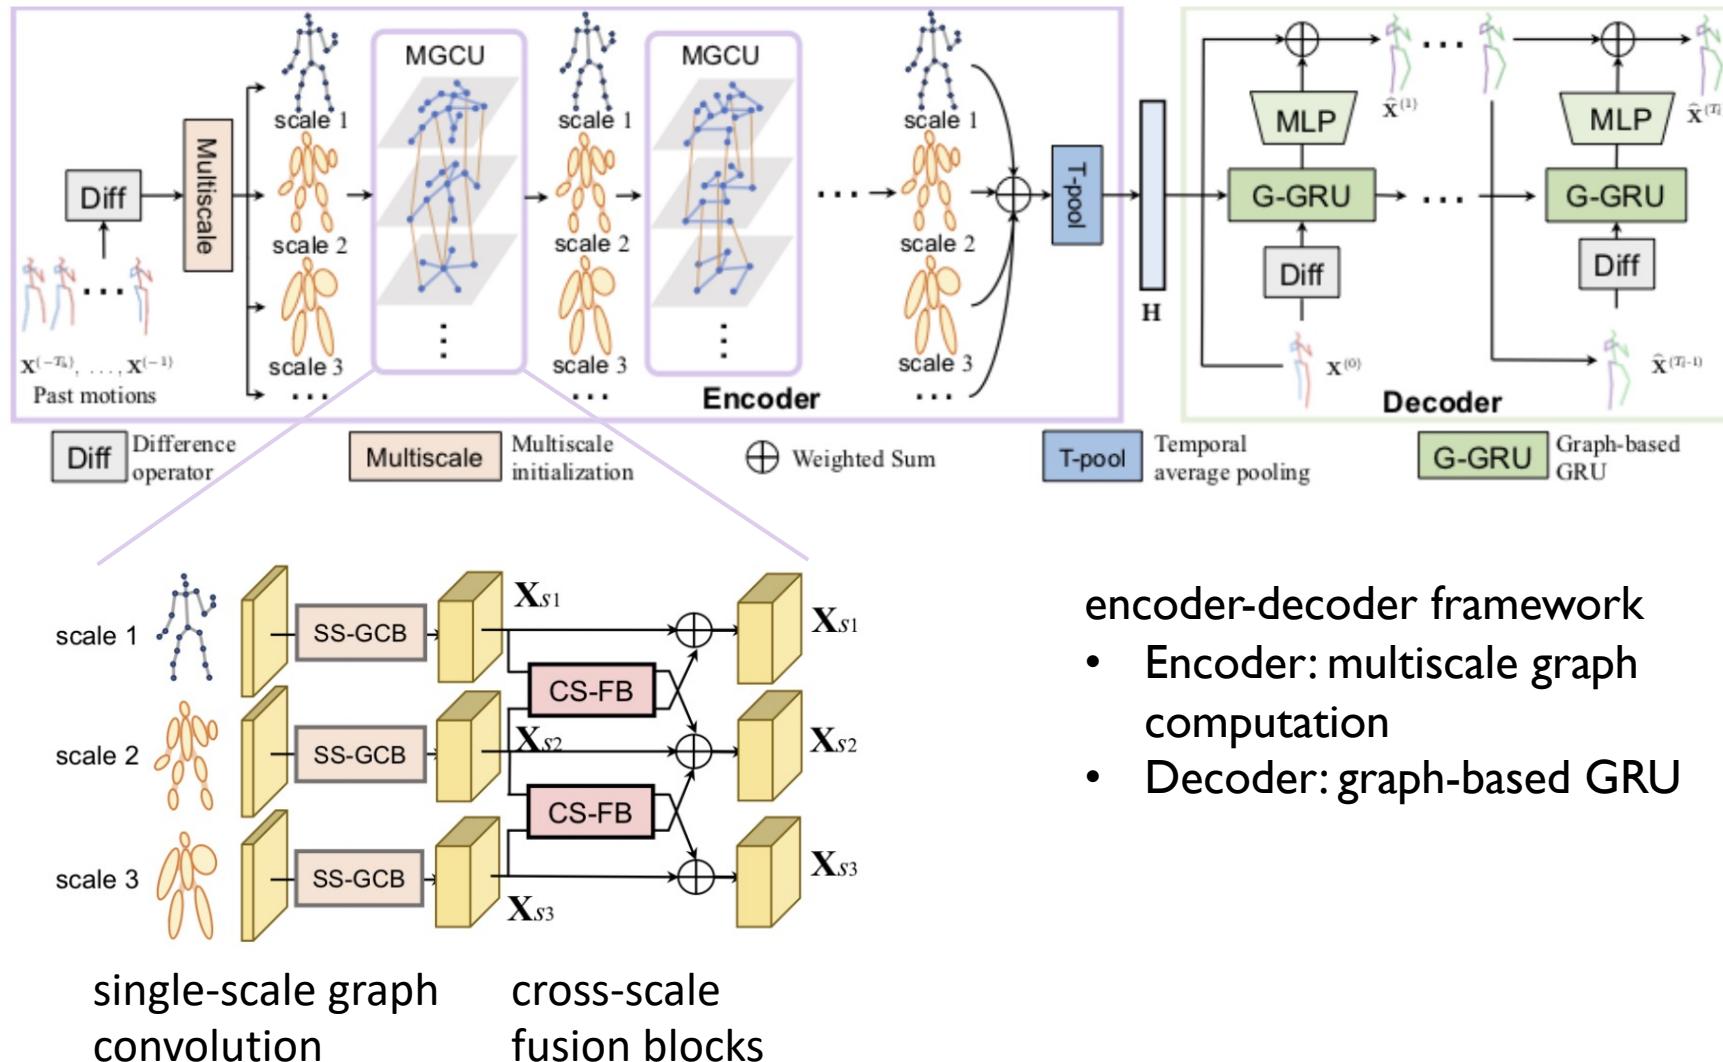
**Capture relations across scales**



**Examples:** local relations Distant relations

# Multiscale GCN model for Human Motion Prediction

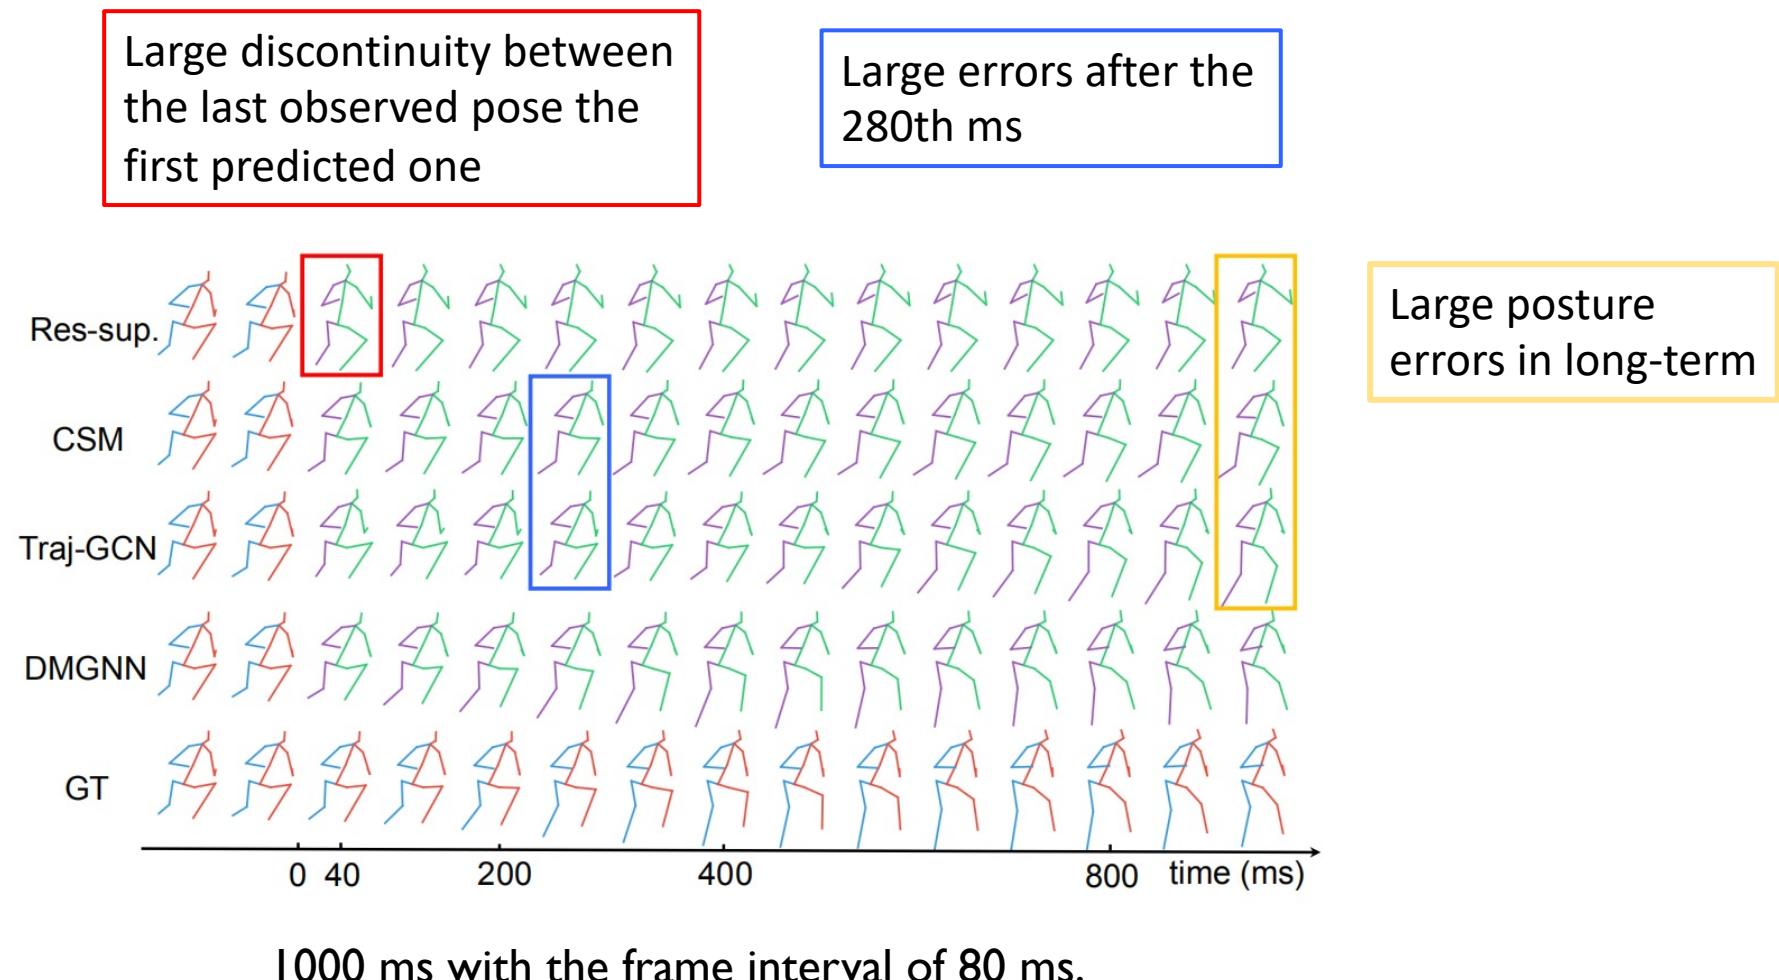
Li, Chen, et al. Dynamic Multiscale Graph Neural Networks for 3D skeleton-based Human Motion Prediction. CVPR 2020 (oral)



# Multiscale GCN model for Human Motion Prediction

Li, Chen, et al. Dynamic Multiscale Graph Neural Networks for 3D skeleton-based Human Motion Prediction. CVPR 2020 (oral)

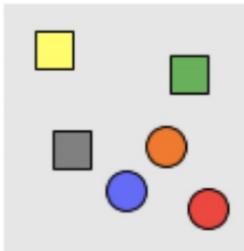
Qualitative comparison on the action ‘Taking Photo’ of H3.6M for both short and long-term prediction.



# Conclusions

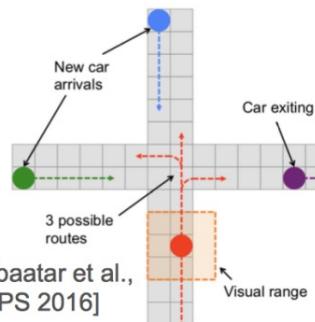
- Deep learning on graphs works and is very effective!
- Exciting area: lots of new applications and extensions (hard to keep up)

Relational reasoning



[Santoro et al., NIPS 2017]

Multi-Agent RL



[Sukhbaatar et al.,  
NIPS 2016]

GCN for recommendation on 16 billion edge graph!



Source pin

[Leskovec lab, Stanford]



SUCCESSFUL  
RECOMMENDATION

BAD RECOMMENDATION

## Open problems:

- Theory
- Scalable, stable generative models
- Learning on large, evolving data
- Multi-modal and cross-model learning (e.g., sequence2graph)

# Next lecture: Motion & tracking



Thank you!

sihengc@sjtu.edu.cn